



# BEA WebLogic Enterprise

## Using Transactions

WebLogic Enterprise 5.1  
Document Edition 5.1  
May 2000

## Copyright

Copyright © 2000 BEA Systems, Inc. All Rights Reserved.

## Restricted Rights Legend

This software and documentation is subject to and made available only pursuant to the terms of the BEA Systems License Agreement and may be used or copied only in accordance with the terms of that agreement. It is against the law to copy the software except as specifically allowed in the agreement. This document may not, in whole or in part, be copied, photocopied, reproduced, translated, or reduced to any electronic medium or machine-readable form without prior consent, in writing, from BEA Systems, Inc.

Use, duplication or disclosure by the U.S. Government is subject to restrictions set forth in the BEA Systems License Agreement and in subparagraph (c)(1) of the Commercial Computer Software-Restricted Rights Clause at FAR 52.227-19; subparagraph (c)(1)(ii) of the Rights in Technical Data and Computer Software clause at DFARS 252.227-7013, subparagraph (d) of the Commercial Computer Software--Licensing clause at NASA FAR supplement 16-52.227-86; or their equivalent.

Information in this document is subject to change without notice and does not represent a commitment on the part of BEA Systems, Inc. THE SOFTWARE AND DOCUMENTATION ARE PROVIDED "AS IS" WITHOUT WARRANTY OF ANY KIND INCLUDING WITHOUT LIMITATION, ANY WARRANTY OF MERCHANTABILITY OR FITNESS FOR A PARTICULAR PURPOSE. FURTHER, BEA Systems, Inc. DOES NOT WARRANT, GUARANTEE, OR MAKE ANY REPRESENTATIONS REGARDING THE USE, OR THE RESULTS OF THE USE, OF THE SOFTWARE OR WRITTEN MATERIAL IN TERMS OF CORRECTNESS, ACCURACY, RELIABILITY, OR OTHERWISE.

## Trademarks or Service Marks

BEA, BEA Builder, BEA Jolt, BEA Manager, BEA MessageQ, BEA Tuxedo, BEA TOP END, BEA WebLogic, and ObjectBroker are registered trademarks of BEA Systems, Inc. BEA eLink, BEA eSolutions, BEA TAP, BEA WebLogic Commerce Server, BEA WebLogic Enterprise, BEA WebLogic Express, BEA WebLogic Personalization Server, BEA WebLogic Server, Java Enterprise Tuxedo, and WebLogic Enterprise Connectivity are trademarks of BEA Systems, Inc.

All other company names may be trademarks of the respective companies with which they are associated.

## Using Transactions

Document Edition	Date	Software Version
5.1	May 2000	BEA WebLogic Enterprise 5.1

---

# Contents

## About This Document

What You Need to Know .....	x
e-docs Web Site .....	x
How to Print the Document.....	x
Related Information.....	xi
Contact Us! .....	xi
Documentation Conventions .....	xii

## 1. Introducing Transactions

Overview of Transactions in WebLogic Enterprise Applications.....	1-2
ACID Properties of Transactions .....	1-2
Supported Programming Models .....	1-2
Supported API Models .....	1-3
Support for Business Transactions .....	1-4
Distributed Transactions and the Two-Phase Commit Protocol .....	1-4
When to Use Transactions.....	1-5
What Happens During a Transaction .....	1-6
Transactions in WebLogic Enterprise CORBA Applications.....	1-6
Transactions in WebLogic Enterprise EJB Applications .....	1-8
Transactions in WebLogic Enterprise RMI Applications .....	1-11
Transactions Sample Code .....	1-12
Transactions Sample CORBA Application .....	1-13
Transactions Sample EJB Code .....	1-25
Transactions Sample RMI Code .....	1-27

## 2. Transaction Service

About the Transaction Service .....	2-2
-------------------------------------	-----

---

Capabilities and Limitations .....	2-2
Lightweight Clients with Delegated Commit .....	2-3
Transaction Propagation (CORBA Only) .....	2-4
Transaction Integrity .....	2-4
Transaction Termination .....	2-4
Flat Transactions .....	2-5
Interoperability Between Remote Clients and the WebLogic Enterprise Domain .....	2-5
Intradomain and Interdomain Interoperability .....	2-5
Network Interoperability .....	2-6
Relationship of the Transaction Service to Transaction Processing .....	2-6
Process Failure .....	2-7
Multithreaded Transaction Client Support .....	2-7
General Constraints .....	2-7
Transaction Service in CORBA Applications .....	2-8
Getting Initial References to the TransactionCurrent Object .....	2-9
CORBA Transaction Service API .....	2-9
CORBA Transaction Service API Extensions .....	2-19
Notes on Using Transactions in WebLogic Enterprise CORBA Applications 2-20	
Transaction Service in EJB Applications .....	2-23
Transaction Service in RMI Applications .....	2-23
UserTransaction API .....	2-24
UserTransaction Methods .....	2-24
Exceptions Thrown by UserTransaction Methods .....	2-26

### 3. Transactions in CORBA Server Applications

Integrating Transactions in a WebLogic Enterprise Client and Server Application 3-2	
Transaction Support in CORBA Applications .....	3-2
Making an Object Automatically Transactional .....	3-4
Enabling an Object to Participate in a Transaction .....	3-5
Preventing an Object from Being Invoked While a Transaction Is Scoped .. 3-6	
Excluding an Object from an Ongoing Transaction .....	3-7
Assigning Policies .....	3-7

---

Using an XA Resource Manager.....	3-8
Opening an XA Resource Manager .....	3-8
Closing an XA Resource Manager .....	3-10
Transactions and Object State Management .....	3-10
Delegating Object State Management to an XA Resource Manager .....	3-10
Waiting Until Transaction Work Is Complete Before Writing to the Database	
3-11	
User-defined Exceptions .....	3-13
About User-defined Exceptions .....	3-13
Defining the Exception.....	3-13
Throwing the Exception .....	3-14
How the Transactions University Sample Application Works (C++ Only)....	3-14
About the Transactions University Sample Application.....	3-15
Transactional Model Used by the Transactions University Sample	
Application.....	3-16
Object State Considerations for the University Server Application .....	3-17
Configuration Requirements for the Transactions Sample Application .	3-19

## 4. Transactions in CORBA Client Applications

Overview of WebLogic Enterprise CORBA Transactions .....	4-2
Summary of the Development Process for Transactions .....	4-2
Step 1: Use the Bootstrap Object to Obtain the TransactionCurrent Object.....	4-3
C++ Example .....	4-3
Java Example.....	4-3
Visual Basic Example .....	4-4
Step 2: Use the TransactionCurrent Methods.....	4-4
C++ Example .....	4-6
Java Example.....	4-6
Visual Basic Example .....	4-7

## 5. Transactions in EJB Applications

Before You Begin .....	5-2
General Guidelines .....	5-2
Transaction Attributes .....	5-3
About Transaction Attributes for EJBs .....	5-3
Transaction Attributes for Container-managed Transactions .....	5-4

---

Transaction Attributes for Bean-managed Transactions .....	5-5
Participating in a Transaction .....	5-5
Transaction Semantics .....	5-5
Transaction Semantics for Container-managed Transactions .....	5-6
Transaction Semantics for Bean-managed Transactions.....	5-9
Session Synchronization .....	5-10
Setting Transaction Timeouts .....	5-11
Handling Exceptions in EJB Transactions .....	5-11

## **6. Transactions in RMI Applications**

Before You Begin .....	6-2
General Guidelines .....	6-2

## **7. Transactions and the WebLogic Enterprise JDBC/XA Driver**

Before You Begin .....	7-2
About Transactions and the WebLogic Enterprise JDBC/XA Driver.....	7-2
Support for Transactions Using the WebLogic Enterprise JDBC/XA Driver	7-2
Local Versus Distributed (Global) Transactions .....	7-4
Transaction Contexts in WebLogic Enterprise JDBC/XA Connections ....	7-7
JDBC Accessibility in Java Methods .....	7-8
JDBC/XA Accessibility in CORBA Methods.....	7-9
JDBC/XA Accessibility in EJB Methods.....	7-9
Using the JDBC/XA Driver.....	7-12
Implementing Distributed Transactions .....	7-14
Importing Packages .....	7-15
Initializing the TransactionCurrent Object Reference.....	7-16
Finding the Connection Pool via JNDI .....	7-16
Setting Up XA Distributed Transactions.....	7-17
Performing a Distributed Transaction .....	7-17

## **8. Administering Transactions**

Modifying the UBBCONFIG File to Accommodate Transactions .....	8-2
Summary of Steps.....	8-2
Step 1: Specify Application-wide Transactions in the RESOURCES Section	8-3

---

Step 2: Create a Transaction Log (TLOG).....	8-3
Step 3: Define Each Resource Manager (RM) and the Transaction Manager Server in the GROUPS Section .....	8-6
Step 4: Enable an Interface to Begin a Transaction .....	8-8
Modifying the Domain Configuration File to Support Transactions (WebLogic Enterprise Servers) .....	8-11
Characteristics of the DMTLOGDEV, DMTLOGNAME, DMTLOGSIZE, MAXRDTRAN, and MAXTRAN Parameters .....	8-11
Characteristics of the AUTOTRAN and TRANTIME Parameters (WebLogic Enterprise CORBA and Tuxedo Servers) .....	8-13
Sample Distributed Application Using Transactions .....	8-14
RESOURCES Section.....	8-14
MACHINES Section .....	8-15
GROUPS and NETWORK Sections.....	8-16
SERVERS, SERVICES, and ROUTING Sections .....	8-17





---

# About This Document

This document explains how to use transactions in CORBA, EJB, and RMI applications that run in the BEA WebLogic Enterprise™ environment.

This document covers the following topics:

- Chapter 1, “Introducing Transactions,” introduces transactions in CORBA, EJB, and RMI applications running in the WebLogic Enterprise environment.
- Chapter 2, “Transaction Service,” describes the WebLogic Enterprise Transaction Service.
- Chapter 3, “Transactions in CORBA Server Applications,” describes how to implement transactions in CORBA C++ and Java server applications.
- Chapter 4, “Transactions in CORBA Client Applications,” describes how to implement transactions in CORBA client applications.
- Chapter 5, “Transactions in EJB Applications,” describes how to implement transactions in EJB applications.
- Chapter 6, “Transactions in RMI Applications,” describes how to implement transactions in RMI applications.
- Chapter 7, “Transactions and the WebLogic Enterprise JDBC/XA Driver,” describes how to use the WebLogic Enterprise JDBC/XA driver in conjunctions with distributed transactions in WebLogic Enterprise Java applications.
- Chapter 8, “Administering Transactions,” describes how to administer transactions in the WebLogic Enterprise environment.

---

# What You Need to Know

This document is intended primarily for application developers who are interested in building transactional C++ and Java applications that run in the WebLogic Enterprise environment. It assumes a familiarity with the WebLogic Enterprise platform, C++ or Java programming, and transaction processing concepts.

## e-docs Web Site

The BEA WebLogic Enterprise product documentation is available on the BEA Systems, Inc. corporate Web site. From the BEA Home page, click the Product Documentation button or go directly to the “e-docs” Product Documentation page at <http://e-docs.bea.com>.

## How to Print the Document

You can print a copy of this document from a Web browser, one file at a time, by using the File—>Print option on your Web browser.

A PDF version of this document is available on the WebLogic Enterprise documentation Home page on the e-docs Web site (and also on the documentation CD). You can open the PDF in Adobe Acrobat Reader and print the entire document (or a portion of it) in book format. To access the PDFs, open the WebLogic Enterprise documentation Home page, click the PDF Files button, and select the document you want to print.

If you do not have the Adobe Acrobat Reader installed, you can download it for free from the Adobe Web site at <http://www.adobe.com/>.

## Related Information

For more information about CORBA, Java 2 Enterprise Edition (J2EE), BEA Tuxedo®, distributed object computing, transaction processing, C++ programming, and Java programming, see the *WebLogic Enterprise Bibliography* in the WebLogic Enterprise online documentation.

## Contact Us!

Your feedback on the BEA WebLogic Enterprise documentation is important to us. Send us e-mail at **docsupport@bea.com** if you have questions or comments. Your comments will be reviewed directly by the BEA professionals who create and update the WebLogic Enterprise documentation.

In your e-mail message, please indicate that you are using the documentation for the BEA WebLogic Enterprise 5.1 release.

If you have any questions about this version of BEA WebLogic Enterprise, or if you have problems installing and running BEA WebLogic Enterprise, contact BEA Customer Support through BEA WebSUPPORT at [www.bea.com](http://www.bea.com). You can also contact Customer Support by using the contact information provided on the Customer Support Card, which is included in the product package.

When contacting Customer Support, be prepared to provide the following information:

- Your name, e-mail address, phone number, and fax number
- Your company name and company address
- Your machine type and authorization codes
- The name and version of the product you are using
- A description of the problem and the content of pertinent error messages

---

# Documentation Conventions

The following documentation conventions are used throughout this document.

Convention	Item
<b>boldface text</b>	Indicates terms defined in the glossary.
Ctrl+Tab	Indicates that you must press two or more keys simultaneously.
<i>italics</i>	Indicates emphasis or book titles.
monospace text	Indicates code samples, commands and their options, data structures and their members, data types, directories, and filenames and their extensions. Monospace text also indicates text that you must enter from the keyboard.  <i>Examples:</i>  #include <iostream.h> void main ( ) the pointer psz chmod u+w * \\tux\data\ap .doc tux.doc BITMAP float
<b>monospace boldface text</b>	Identifies significant words in code.  <i>Example:</i>  void <b>commit</b> ( )
<i>monospace italic text</i>	Identifies variables in code.  <i>Example:</i>  String <i>expr</i>
UPPERCASE TEXT	Indicates device names, environment variables, and logical operators.  <i>Examples:</i>  LPT1 SIGNON OR

Convention	Item
{ }	Indicates a set of choices in a syntax line. The braces themselves should never be typed.
[ ]	Indicates optional items in a syntax line. The brackets themselves should never be typed. <i>Example:</i> <code>buildobjclient [-v] [-o name ] [-f file-list]... [-l file-list]...</code>
	Separates mutually exclusive choices in a syntax line. The symbol itself should never be typed.
...	Indicates one of the following in a command line: <ul style="list-style-type: none"><li>■ That an argument can be repeated several times in a command line</li><li>■ That the statement omits additional optional arguments</li><li>■ That you can enter additional parameters, values, or other information</li></ul> The ellipsis itself should never be typed. <i>Example:</i> <code>buildobjclient [-v] [-o name ] [-f file-list]... [-l file-list]...</code>
. . . .	Indicates the omission of items from a code example or from a syntax line. The vertical ellipsis itself should never be typed.



# 1 Introducing Transactions

This topic includes the following sections:

- Overview of Transactions in WebLogic Enterprise Applications
- When to Use Transactions
- What Happens During a Transaction
- Transactions Sample Code

# Overview of Transactions in WebLogic Enterprise Applications

This topic includes the following sections:

- ACID Properties of Transactions
- Supported Programming Models
- Supported API Models
- Support for Business Transactions
- Distributed Transactions and the Two-Phase Commit Protocol

## ACID Properties of Transactions

One of the most fundamental features of the WebLogic Enterprise system is transaction management. Transactions are a means to guarantee that database transactions are completed accurately and that they take on all the ACID properties (atomicity, consistency, isolation, and durability) of a high-performance transaction. WebLogic Enterprise protects the integrity of your transactions by providing a complete infrastructure for ensuring that database updates are done accurately, even across a variety of resource managers (RMs). If any one of the operations fails, the entire set of operations is rolled back.

## Supported Programming Models

WebLogic Enterprise supports transactions in two different programming models:

- The Object Management Group Common Object Request Broker (CORBA) in C++ and Java applications, in compliance with the *The Common Object Request Broker: Architecture and Specification*, Revision 2.2, February 1998.



- The Sun Microsystems, Inc., Java 2 Platform, Enterprise Edition (J2EE). WebLogic Enterprise provides full support for transactions in Java applications that use Enterprise JavaBeans, in compliance with the Enterprise JavaBeans Specification 1.1, published by Sun Microsystems, Inc. WebLogic Enterprise also supports the Java Transaction API (JTA) Specification version 1.0.1, also published by Sun Microsystems, Inc.

## Supported API Models

WebLogic Enterprise supports two transaction API models:

- CORBA services Object Transaction Service (OTS) and the Java Transaction Service (JTS).

WebLogic Enterprise provides a C++ interface to the OTS and a Java interface to the OTS and the JTS. The JTS is the Sun Microsystems, Inc. Java interface for transaction services, and is based on the OTS. The OTS and the JTS are accessed through the `org.omg.CosTransactions.Current` environmental object. For information about using the `TransactionCurrent` environmental object, see the *C++ Bootstrap Object Programming Reference* or the *Java Bootstrap Object Programming Reference*.

- The Sun Microsystems, Inc. Java Transaction API (JTA), which is used by:
  - CORBA applications within BEA's TP Framework.
  - Enterprise JavaBean (EJB) applications within the WebLogic Enterprise EJB container.
  - Remote Method Invocation (RMI) applications within the WebLogic Enterprise infrastructure.

Only the application-level demarcation interface (`javax.transaction.UserTransaction`) is supported. For information about JTA, see the following sources:

- The `javax.transaction` package description in the *WebLogic Enterprise Javadoc*.
- The Java Transaction API specification, published by Sun Microsystems, Inc. and available from the Sun Microsystems, Inc. Web site ([www.sun.com](http://www.sun.com)).

# Support for Business Transactions

OTS, JTS, and JTA each provide the following support for your business transactions:

- Creates a global transaction identifier when a client application initiates a transaction.
- Works with the WebLogic Enterprise infrastructure to track objects that are involved in a transaction and, therefore, need to be coordinated when the transaction is ready to commit.
- Notifies the resource managers—which are, most often, databases—when they are accessed on behalf of a transaction. Resource managers then lock the accessed records until the end of the transaction.
- Orchestrates the two-phase commit when the transaction completes, which ensures that all the participants in the transaction commit their updates simultaneously. It coordinates the commit with any databases that are being updated using Open Group's XA protocol. Almost all relational databases support this standard.
- Executes the rollback procedure when the transaction must be stopped.
- Executes a recovery procedure when failures occur. It determines which transactions were active in the machine at the time of the crash, and then determines whether the transaction should be rolled back or committed.

## Distributed Transactions and the Two-Phase Commit Protocol

WebLogic Enterprise supports distributed transactions and the two-phase commit protocol for enterprise applications. A *distributed transaction* is a transaction that updates multiple resource managers (such as databases) in a coordinated manner. The *two-phase commit protocol (2PC)* is a method of coordinating a single transaction across one or more resource managers. It guarantees data integrity by ensuring that transactional updates are committed in all of the participating databases, or are fully rolled back out of all the databases, reverting to the state prior to the start of the transaction.

# When to Use Transactions

Transactions are appropriate in the situations described in the following list. Each situation describes a transaction model supported by the WebLogic Enterprise system.

- The client application needs to make invocations on several objects, which may involve write operations to one or more databases. If any one invocation is unsuccessful, any state that is written (either in memory or, more typically, to a database) must be rolled back.

For example, consider a travel agent application. The client application needs to arrange for a journey to a distant location; for example, from Strasbourg, France, to Alice Springs, Australia. Such a journey would inevitably require multiple individual flight reservations. The client application works by reserving each individual segment of the journey in sequential order; for example, Strasbourg to Paris, Paris to New York, New York to Los Angeles. However, if any individual flight reservation cannot be made, the client application needs a way to cancel all the flight reservations made up to that point.

- The client application needs a conversation with an object managed by the server application, and the client application needs to make multiple invocations on a specific object instance. The conversation may be characterized by one or more of the following:
  - Data is cached in memory or written to a database during or after each successive invocation.
  - Data is written to a database at the end of the conversation.
  - The client application needs the object to maintain an in-memory context between each invocation; that is, each successive invocation uses the data that is being maintained in memory across the conversation.
  - At the end of the conversation, the client application needs the ability to cancel all database write operations that may have occurred during or at the end of the conversation.

For example, consider an Internet-based online shopping cart application. Users of the client application browse through an online catalog and make multiple purchase selections. When the users are done choosing all the items they want to buy, they proceed to check out and enter their credit card information to make the purchase. If the credit card check fails, the shopping application needs a way

to cancel all the pending purchase selections in the shopping cart, or roll back any purchase transactions made during the conversation.

- Within the scope of a single client invocation on an object, the object performs multiple edits to data in a database. If one of the edits fails, the object needs a mechanism to roll back all the edits. (In this situation, the individual database edits are not necessarily CORBA, EJB, or RMI invocations.)

For example, consider a banking application. The client invokes the transfer operation on a teller object. The transfer operation requires the teller object to make the following invocations on the bank database:

- Invoking the debit method on one account.
- Invoking the credit method on another account.

If the credit invocation on the bank database fails, the banking application needs a way to roll back the previous debit invocation.

# What Happens During a Transaction


This topic includes the following sections:

- Transactions in WebLogic Enterprise CORBA Applications
- Transactions in WebLogic Enterprise EJB Applications
- Transactions in WebLogic Enterprise RMI Applications

## Transactions in WebLogic Enterprise CORBA Applications

Figure 1-1 illustrates how transactions work in a WebLogic Enterprise CORBA application.

**Figure 1-1 How Transactions Work in a WebLogic Enterprise CORBA Application**



For CORBA applications, a basic transaction works in the following way:

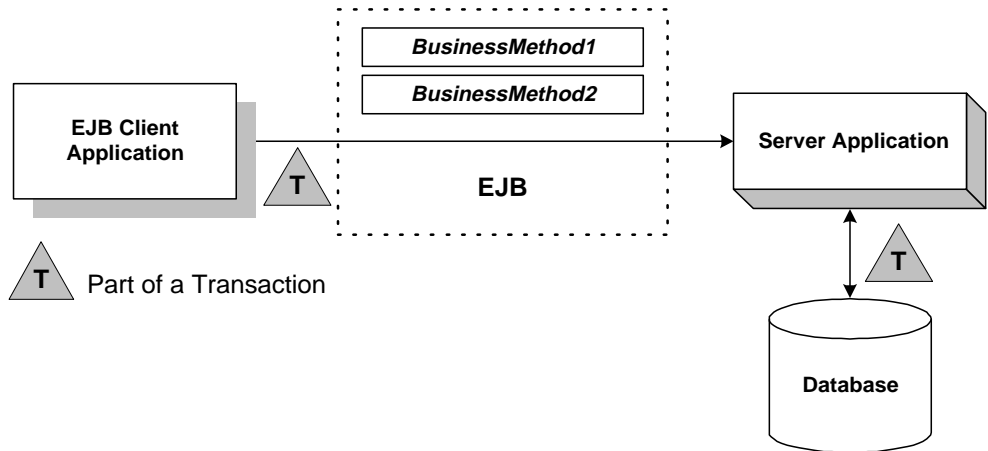
1. The client application uses the `Bootstrap` object to return an object reference to the `TransactionCurrent` object for the WebLogic Enterprise domain.
2. A client application begins a transaction using the `Tobj::TransactionCurrent::begin()` operation, and issues a request to the CORBA interface through the TP Framework. All operations on the CORBA interface execute within the scope of a transaction.
  - If a call to any of these operations raises an exception (either explicitly or as a result of a communication failure), the exception can be caught and the transaction can be rolled back.
  - If no exceptions occur, the client application commits the current transaction using the `Tobj::TransactionCurrent::commit()` operation. This operation ends the transaction and starts the processing of the operation. The transaction is committed only if all of the participants in the transaction agree to commit.

3. The `Tobj::TransactionCurrent::commit()` operation causes the TP Framework to call the transaction manager to complete the transaction.
4. The transaction manager is responsible for coordinating with the resource managers to update the database.

# Transactions in WebLogic Enterprise EJB Applications

Figure 1-2 illustrates how transactions work in a WebLogic Enterprise EJB application.

**Figure 1-2 How Transactions Work in a WebLogic Enterprise EJB Application**



WebLogic Enterprise supports two types of transactions in WebLogic Enterprise EJB applications:

- In *container-managed transactions*, the WebLogic Enterprise EJB container manages the transaction demarcation. Transaction attributes in the EJB deployment descriptor determine how the WebLogic Enterprise EJB container handles transactions with each method invocation. For more information about the deployment descriptor, see the *WebLogic EJB XML Reference*.
- In *bean-managed transactions*, the EJB manages the transaction demarcation. The EJB makes explicit method invocations on the `UserTransaction` object to

begin, commit, and roll back transactions. For more information about the `UserTransaction` object, see “UserTransaction API” on page 2-24.

The sequence of transaction events differs between container-managed and bean-managed transactions.

### Container-managed Transactions

For EJB applications with container-managed transactions, a basic transaction works in the following way:

1. In the EJB’s deployment descriptor, the Bean Provider or Application Assembler specifies the transaction type (`transaction-type` element) for container-managed demarcation (`Container`).
2. In the EJB’s deployment descriptor, the Bean Provider or Application Assembler specifies the default transaction attribute (`trans-attribute` element) for the EJB, which is one of the following settings: `NotSupported`, `Required`, `Supports`, `RequiresNew`, `Mandatory`, or `Never`. For a detailed description of these settings, see Section 11.6.2 in the Enterprise JavaBeans Specification 1.1, published by Sun Microsystems, Inc.
3. Optionally, in the EJB’s deployment descriptor, the Bean Provider or Application Assembler specifies the `trans-attribute` for one or more methods.
4. When a client application invokes a method in the EJB, the EJB container checks the `trans-attribute` setting in the deployment descriptor for that method. If no setting is specified for the method, the EJB uses the default `trans-attribute` setting for that EJB.
5. The EJB container takes the appropriate action depending on the applicable `trans-attribute` setting.
  - For example, if the `trans-attribute` setting is `Required`, the EJB container invokes the method within the existing transaction context or, if the client called without a transaction context, the EJB container begins a new transaction before executing the method.
  - In another example, if the `trans-attribute` setting is `Mandatory`, the EJB container invokes the method within the existing transaction context. If the client called without a transaction context, the EJB container throws the `javax.transaction.TransactionRequiredException` exception.

6. During invocation of the business method, if it is determined that a rollback is required, the business method calls the `EJBContext.setRollbackOnly` method, which notifies the EJB container that the transaction is to be rolled back at the end of the method invocation.

**Note:** Calling the `EJBContext.setRollbackOnly` method is allowed only for methods that have a meaningful transaction context.

7. At the end of the method execution and before the result is sent to the client, the EJB container completes the transaction, either by committing the transaction or rolling it back (if the `EJBContext.setRollbackOnly` method was called).

You can control transaction timeouts by setting the `trans-timeout-seconds` element in the `weblogic-ejb-extensions.xml` file. For more information about the `weblogic-ejb-extensions.xml` file, see the WebLogic Enterprise *EJB XML Reference*. You can also change this setting with the WebLogic Enterprise EJB Deployer, as described in *Using the WebLogic Enterprise EJB Deployer*.

## Bean-managed Transactions

For EJB applications with bean-managed transaction demarcations, a basic transaction works in the following way:

1. In the EJB's deployment descriptor, the Bean Provider or Application Assembler specifies the transaction type (`transaction-type` element) for container-managed demarcation (Bean).
2. The client application uses JNDI to obtain an object reference to the `UserTransaction` object for the WebLogic Enterprise domain.
3. The client application begins a transaction using the `UserTransaction.begin` method, and issues a request to the EJB through the EJB container. All operations on the EJB execute within the scope of a transaction.
  - If a call to any of these operations raises an exception (either explicitly or as a result of a communication failure), the exception can be caught and the transaction can be rolled back using the `UserTransaction.rollback` method.
  - If no exceptions occur, the client application commits the current transaction using the `UserTransaction.commit` method. This method ends the transaction and starts the processing of the operation. The transaction is committed only if all of the participants in the transaction agree to commit.

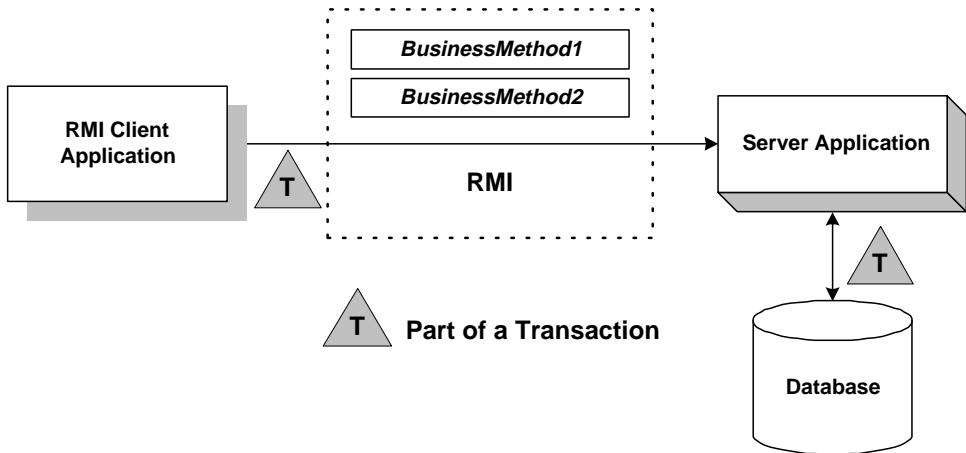


4. The `UserTransaction.commit` method causes the EJB container to call the transaction manager to complete the transaction.
5. The transaction manager is responsible for coordinating with the resource managers to update any databases.

## Transactions in WebLogic Enterprise RMI Applications

Figure 1-3 illustrates how transactions work in a WebLogic Enterprise RMI application.

**Figure 1-3 How Transactions Work in a WebLogic Enterprise RMI Application**



For RMI client and server applications, a basic transaction works in the following way:

1. The application uses JNDI to return an object reference to the `UserTransaction` object for the WebLogic Enterprise domain.

Obtaining the object reference begins a conversational state between the application and that object. The conversational state continues until the transaction is completed (committed or rolled back). Once instantiated, RMI objects remain active in memory until they are released (typically during server shutdown). For the duration of the transaction, the WebLogic Enterprise infrastructure does not perform any deactivation or activation.

2. The client application begins a transaction using the `UserTransaction.begin` method, and issues a request to the server application. All operations on the server application execute within the scope of a transaction.
  - If a call to any of these operations raises an exception (either explicitly or as a result of a communication failure), the exception can be caught and the transaction can be rolled back using the `UserTransaction.rollback` method.
  - If no exceptions occur, the client application commits the current transaction using the `UserTransaction.commit` method. This method ends the transaction and starts the processing of the operation. The transaction is committed only if all of the participants in the transaction agree to commit.
3. The `UserTransaction.commit` method causes WebLogic Enterprise to call the transaction manager to complete the transaction.
4. The transaction manager is responsible for coordinating with the resource managers to update any databases.

For guidelines about using transactions in RMI applications, see Chapter 6, “Transactions in RMI Applications.”

# Transactions Sample Code

This topic includes the following sections:

- Transactions Sample CORBA Application
- Transactions Sample EJB Code
- Transactions Sample RMI Code

# Transactions Sample CORBA Application

In the Transactions sample CORBA application, the operation of registering for courses is executed within the scope of a transaction. The transaction model used in the Transactions sample application is a combination of the conversational model and the model in which a single client invocation makes multiple individual operations on a database.

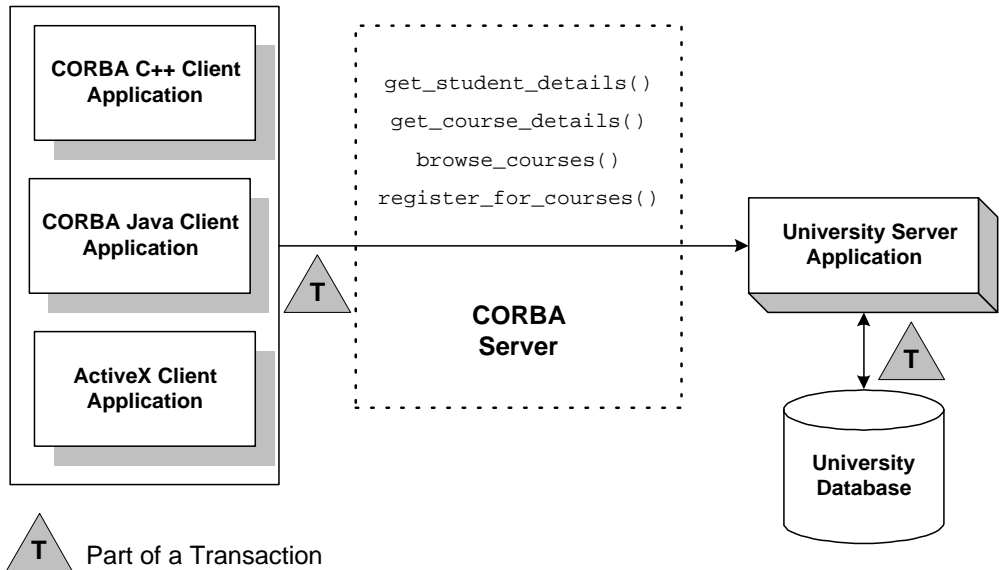
## Workflow for the Transactions Sample Application

The Transactions sample application works in the following way:

1. Students submit a list of courses for which they want to be registered.
2. For each course in the list, the server application checks whether:
  - The course is in the database.
  - The student is already registered for a course.
  - The student exceeds the maximum number of credits the student can take.
3. One of the following occurs:
  - If the course meets all the criteria, the server application registers the student for the course.
  - If the course is not in the database or if the student is already registered for the course, the server application adds the course to a list of courses for which the student could not be registered. After processing all the registration requests, the server application returns the list of courses for which registration failed. The client application can then choose to either commit the transaction (thereby registering the student for the courses for which registration request succeeded) or to roll back the transaction (thus, not registering the student for any of the courses).
  - If the student exceeds the maximum number of credits the student can take, the server application returns a `TooManyCredits` user exception to the client application. The client application provides a brief message explaining that the request was rejected. The client application then rolls back the transaction.

Figure 1-4 illustrates how the Transactions sample application works.

Figure 1-4 Transactions Sample Application



The Transactions sample application shows two ways in which a transaction can be rolled back:

- **Nonfatal.** If the registration for a course fails because the course is not in the database, or because the student is already registered for the course, the server application returns the numbers of those courses to the client application. The decision to roll back the transaction lies with the user of the client application (and the Transaction client application code rolls back the transaction automatically in this case).
- **Fatal.** If the registration for a course fails because the student exceeds the maximum number of credits he or she can take, the server application generates a CORBA exception and returns it to the client. The decision to roll back the transaction also lies with the client application.

Thus, the Transactions sample application also shows how to implement user-defined CORBA exceptions. For example, if the student tries to register for a course that would exceed the maximum number of courses for which the student can register, the server application returns the `TooManyCredits` exception. When the client application receives this exception, the client application rolls back the transaction automatically.

**Note:** For information about how transactions are implemented in CORBA/Java WebLogic Enterprise applications, see the *Transactions Sample* in the WebLogic Enterprise online documentation.

## Development Steps

This topic describes the following development steps for writing a WebLogic Enterprise application that contains transaction processing code:

- Step 1: Writing the OMG IDL
- Step 2: Defining Transaction Policies for the Interfaces
- Step 3: Writing the Server Application
- Step 4: Writing the Client Application
- Step 5: Creating a Configuration File

The Transactions sample application is used to demonstrate these development steps. The source files for the Transactions sample application are located in the `\samples\corba\university` directory of the WebLogic Enterprise software. For information about building and running the Transactions sample application, see the *Transactions Sample* in the WebLogic Enterprise online documentation.

The XA Bankapp sample application demonstrates how to use transactions in Java WebLogic Enterprise applications. The source files for the XA Bankapp sample application are located in the `\samples\corba\bankapp_java\XA` directory of the WebLogic Enterprise software. For information about building and running the XA Bankapp sample application, see the *Bankapp Sample Using XA* in the WebLogic Enterprise online documentation.

### Step 1: Writing the OMG IDL

You need to specify interfaces involved in transactions in Object Management Group (OMG) Interface Definition Language (IDL) just as you would any other CORBA interface. You must also specify any user exceptions that might occur from using the interface.

For the Transactions sample application, you would define in OMG IDL the `Registrar` interface and the `register_for_courses()` operation. The `register_for_courses()` operation has a parameter, `NotRegisteredList`, which

returns to the client application the list of courses for which registration failed. If the value of `NotRegisteredList` is empty, then the client application commits the transaction. You also need to define the `TooManyCredits` user exception.

Listing 1-1 includes the OMG IDL for the Transactions sample application.

## **Listing 1-1   OMG IDL for the Transactions Sample Application**

---

```
#pragma prefix "beasys.com"
module UniversityT
{
    typedef unsigned long CourseNumber;
    typedef sequence<CourseNumber> CourseNumberList;

    struct CourseSynopsis
    {
        CourseNumber    course_number;
        string           title;
    };
    typedef sequence<CourseSynopsis> CourseSynopsisList;

    interface CourseSynopsisEnumerator
    {
        //Returns a list of length 0 if there are no more entries
        CourseSynopsisList get_next_n(
            in unsigned long number_to_get, // 0 = return all
            out unsigned long number_remaining
        );

        void destroy();
    };
    typedef unsigned short Days;
    const Days MONDAY    = 1;
    const Days TUESDAY   = 2;
    const Days WEDNESDAY = 4;
    const Days THURSDAY  = 8;
    const Days FRIDAY    = 16;
}
//Classes restricted to same time block on all scheduled days,
//starting on the hour

struct ClassSchedule
{
    Days            class_days; // bitmask of days
    unsigned short start_hour; // whole hours in military time
}
```

```
        unsigned short duration;    // minutes
};

struct CourseDetails
{
    CourseNumber    course_number;
    double          cost;
    unsigned short  number_of_credits;
    ClassSchedule   class_schedule;
    unsigned short  number_of_seats;
    string          title;
    string          professor;
    string          description;
};

typedef sequence<CourseDetails> CourseDetailsList;
typedef unsigned long StudentId;

struct StudentDetails
{
    StudentId      student_id;
    string         name;
    CourseDetailsList registered_courses;
};

enum NotRegisteredReason
{
    AlreadyRegistered,
    NoSuchCourse
};

struct NotRegistered
{
    CourseNumber    course_number;
    NotRegisteredReason not_registered_reason;
};

typedef sequence<NotRegistered> NotRegisteredList;

exception TooManyCredits
{
    unsigned short maximum_credits;
};

//The Registrar interface is the main interface that allows
//students to access the database.
interface Registrar
{
    CourseSynopsisList
    get_courses_synopsis(
        in string          search_criteria,
        in unsigned long    number_to_get,
```

# 1 *Introducing Transactions*

---

```
        out unsigned long          number_remaining,
        out CourseSynopsisEnumerator rest
    );

    CourseDetailsList get_courses_details(in CourseNumberList
        courses);
    StudentDetails get_student_details(in StudentId student);
    NotRegisteredList register_for_courses(
        in StudentId          student,
        in CourseNumberList courses
    ) raises (
        TooManyCredits
    );
};

// The RegistrarFactory interface finds Registrar interfaces.

interface RegistrarFactory
{
    Registrar find_registrar(
    );
};
```

## Step 2: Defining Transaction Policies for the Interfaces

Transaction policies are used on a per-interface basis. During design, it is decided which interfaces within a WebLogic Enterprise application will handle transactions. Table 1-1 describes the CORBA transaction policies.

**Table 1-1 CORBA Transaction Policies**

Transaction Policy	Description
always	The interface must always be part of a transaction. If the interface is not part of a transaction, a transaction will be automatically started by the TP Framework.
ignore	The interface is not transactional. However, requests made to this interface within a scope of a transaction are allowed. The AUTOTRAN parameter, specified in the UBBCONFIG file for this interface, is ignored.



**Table 1-1 CORBA Transaction Policies (Continued)**

Transaction Policy	Description
<code>never</code>	The interface is not transactional. Objects created for this interface can never be involved in a transaction. The WebLogic Enterprise system generates an exception ( <code>INVALID_TRANSACTION</code> ) if an interface with this policy is involved in a transaction.
<code>optional</code>	The interface may be transactional. Objects can be involved in a transaction if the request is transactional. This transaction policy is the default.

During development, you decide which interfaces will execute in a transaction by assigning transaction policies in the following ways:

- For C++ server applications in CORBA, you specify transaction policies in the Implementation Configuration File (ICF). A template ICF file is created by the `genicf` command. For more information about the ICFs, see “Implementation Configuration File (ICF)” in the *CORBA C++ Programming Reference*.
- For Java server applications in CORBA, you specify transaction policies in the Server Description File, written in Extensible Markup Language (XML). For more information about Server Description files, see “Server Description File” in the *CORBA Java Programming Reference*.

In the Transactions sample application, the transaction policy of the `Registrar` interface is set to `always`.

### Step 3: Writing the Server Application

When using transactions in server applications, you need to write methods that implement the interface’s operations. In the Transactions sample application, you would write a method implementation for the `register_for_courses()` operation.

If your WebLogic Enterprise application uses a database, you need to include in the server application code that opens and closes an XA Resource Manager. These operations are included in the `Server::initialize()` and `Server::release()` operations of the `Server` object. Listing 1-2 shows the portion of the code for the `Server` object in the Transactions sample application that opens and closes the XA Resource Manager.

**Note:** For a complete example of a C++ server application that implements transactions, see the *Transactions Sample* in the WebLogic Enterprise online documentation. For an example of a Java server application that implements transactions, see *Bankapp Sample Using XA* in the WebLogic Enterprise online documentation.

## **Listing 1-2 C++ Server Object in Transactions Sample Application**

---

```
CORBA::Boolean Server::initialize(int argc, char* argv[])
{
    TRACE_METHOD("Server::initialize");
    try {
        open_database();
        begin_transactional();
        register_fact();
        return CORBA_TRUE;
    }
    catch (CORBA::Exception& e) {
        LOG("CORBA exception : " <<e);
    }
    catch (SamplesDBException& e) {
        LOG("Can't connect to database");
    }
    catch (...) {
        LOG("Unexpected database error : " <<e);
    }
    catch (...) {
        LOG("Unexpected exception");
    }
    cleanup();
    return CORBA_FALSE;
}

void Server::release()
{
    TRACE_METHOD("Server::release");
    cleanup();
}

static void cleanup()
{
    unregister_factory();
    end_transactional();
    close_database();
}

//Utilities to manage transaction resource manager
```

```
CORBA::Boolean s_became_transactional = CORBA_FALSE;
static void begin_transactional()
{
    TP::open_xa_rm();
    s_became_transactional = CORBA_TRUE;
}
static void end_transactional()
{
    if(!s_became_transactional){
        return//cleanup not necessary
    }
    try {
        TP::close_xa_rm ();
    }
    catch (CORBA::Exception& e) {
        LOG("CORBA Exception : " << e);
    }
    catch (...) {
        LOG("unexpected exception");
    }
    s_became_transactional = CORBA_FALSE;
}
```

## Step 4: Writing the Client Application

The client application needs code that performs the following tasks:

1. Obtains a reference to the `TransactionCurrent` object from the `Bootstrap` object.
2. Begins a transaction by invoking the `Tobj::TransactionCurrent::begin()` operation on the `TransactionCurrent` object.
3. Invokes operations on the object. In the Transactions sample application, the client application invokes the `register_for_courses()` operation on the `Registrar` object, passing a list of courses.

Listing 1-3 illustrates the portion of the CORBA C++ client applications in the Transactions sample application that illustrates the development steps for transactions.

For an example of a CORBA Java client application that uses transactions, see *Bankapp Sample Using XA* in the WebLogic Enterprise online documentation. For an example of using transactions in an ActiveX client application, see Chapter 4, “Transactions in CORBA Client Applications.”

## **Listing 1-3 Transactions Code for CORBA C++ Client Applications**

---

```
CORBA::Object_var var_transaction_current_oref =
    Bootstrap.resolve_initial_references("TransactionCurrent");
CosTransactions::Current_var transaction_current_oref=
    CosTransactions::Current::_narrow(var_transaction_current_oref.in());
//Begin the transaction
var_transaction_current_oref->begin();
try {
    //Perform the operation inside the transaction
    pointer_Registar_ref->register_for_courses(student_id, course_number_list);
    ...
    //If operation executes with no errors, commit the transaction:
    CORBA::Boolean report_heuristics = CORBA_TRUE;
    var_transaction_current_ref->commit(report_heuristics);
}
catch (...) {
    //If the operation has problems executing, rollback the
    //transaction. Then throw the original exception again.
    //If the rollback fails, ignore the exception and throw the
    //original exception again.
    try {
        var_transaction_current_ref->rollback();
    }
    catch (...) {
        TP::userlog("rollback failed");
    }
}
throw;
}
```

### Step 5: Creating a Configuration File

You need to add the following information to the configuration file for a transactional WebLogic Enterprise application:

- In the **GROUPS** section:
  - In the **OPENINFO** parameter, include the information needed to open the resource manager for the database. You obtain this information from the product documentation for your database. Note that the default version of the `com.beasys.Tobj.Server.initialize` method automatically opens the resource manager.
  - In the **CLOSEINFO** parameter, include the information needed to close the resource manager for the database. By default, the **CLOSEINFO** parameter is empty.

- Specify the TMSNAME and TMSCOUNT parameters to associate the XA resource manager with a specified server group.
- In the SERVERS section, define a server group that includes both the server application that includes the interface and the server application that manages the database. This server group needs to be specified as transactional.
- Include the pathname to the transaction log (TLOG) in the TLOGDEVICE parameter. For more information about the transaction log, see Chapter 8, “Administering Transactions.”

Listing 1-4 includes the portions of the configuration file that define this information for the Transactions sample application.

#### **Listing 1-4 Configuration File for Transactions Sample Application**

---

```
*RESOURCES
    IPCKEY      55432
    DOMAINID    university
    MASTER      SITE1
    MODEL        SHM
    LDBAL        N
    SECURITY     APP_PW

*MACHINES
    BLOTTO
    LMID = SITE1
    APPDIR = C:\TRANSACTION_SAMPLE
    TUXCONFIG=C:\TRANSACTION_SAMPLE\tuxconfig
    TLOGDEVICE=C:\APP_DIR\TLOG
    TLOGNAME=TLOG
    TUXDIR="C:\WLEdir"
    MAXWSCLIENTS=10

*GROUPS
    SYS_GRP
        LMID      = SITE1
        GRPNO      = 1
    ORA_GRP
        LMID      = SITE1
        GRPNO      = 2

    OPENINFO = "ORACLE_XA:Oracle_XA+SqlNet=ORCL+Acc=P
/scott/tiger+SesTm=100+LogDir=."+MaxCur=5"
    CLOSEINFO = " "
```

# 1 *Introducing Transactions*

---

```
TMSNAME      = "TMS_ORA"
TMSCOUNT     = 2

*SERVERS
  DEFAULT:
  RESTART    = Y
  MAXGEN     = 5

  TMSYSEVT
    SRVGRP   = SYS_GRP
    SRVID    = 1

  TMFFNAME
    SRVGRP   = SYS_GRP
    SRVID    = 2
    CLOPT    = "-A -- -N -M"

  TMFFNAME
    SRVGRP   = SYS_GRP
    SRVID    = 3
    CLOPT    = "-A -- -N"

  TMFFNAME
    SRVGRP   = SYS_GRP
    SRVID    = 4
    CLOPT    = "-A -- -F"

  TMIFRSVR
    SRVGRP   = SYS_GRP
    SRVID    = 5

  UNIVT_SERVER
    SRVGRP   = ORA_GRP
    SRVID    = 1
    RESTART  = N

  ISL
    SRVGRP   = SYS_GRP
    SRVID    = 6
    CLOPT    = -A -- -n //MACHINENAME:2500

*SERVICES
```

For information about the transaction log and defining parameters in the Configuration file, see Chapter 8, “Administering Transactions.”

# Transactions Sample EJB Code

This topic provides a walkthrough of sample code fragments from a class in an EJB application. This topic includes the following sections:

- Importing Packages
- Initializing the UserTransaction Object
- Using JNDI to Return an Object Reference to the UserTransaction Object
- Starting a Transaction
- Completing a Transaction

The code fragments demonstrate using the `UserTransaction` object for *bean-managed* transaction demarcation. The deployment descriptor for this bean specifies the transaction type (`transaction-type` element) for transaction demarcation (`Bean`).

**Note:** These code fragments do not derive from any of the sample applications that ship with WebLogic Enterprise. They merely illustrate the use of the `UserTransaction` object within an EJB application.

## Importing Packages

Listing 1-5 shows importing the necessary packages for transactions, including:

- `javax.transaction.UserTransaction`. For a list of methods associated with this object, see “UserTransaction Methods” on page 2-24.
- System exceptions. For a list of exceptions, see “Exceptions Thrown by UserTransaction Methods” on page 2-26.

### Listing 1-5 Importing Packages

---

```
import javax.naming.*;
import javax.transaction.UserTransaction;
import javax.transaction.SystemException;
import javax.transaction.HeuristicMixedException;
import javax.transaction.HeuristicRollbackException;
import javax.transaction.NotSupportedException;
```

```
import javax.transaction.RollbackException
import javax.transaction.IllegalStateException
import javax.transaction.SecurityException
```

## Initializing the UserTransaction Object

Listing 1-6 shows initializing an instance of the `UserTransaction` object to null.

### **Listing 1-6 Initializing the UserTransaction Object**

---

```
UserTransaction tx = null;
```

## Using JNDI to Return an Object Reference to the UserTransaction Object

Listing 1-7 shows searching the JNDI tree to return an object reference to the `UserTransaction` object for the appropriate WebLogic Enterprise domain.

### **Listing 1-7 Performing a JNDI Lookup**

---

```
try {
    Context ctx = getInitialContext();
    tx = (UserTransaction)ctx.lookup("java:comp/UserTransaction");
}
```

## Starting a Transaction

Listing 1-8 shows starting a transaction by calling the `javax.transaction.UserTransaction.begin` method. Database operations that occur after this method invocation and prior to completing the transaction exist within the scope of this transaction.

### **Listing 1-8 Starting a Transaction**

---

```
tx.begin();
```



## Completing a Transaction

Listing 1-9 shows completing the transaction depending on whether an exception was thrown during any of the database operations that were attempted within the scope of this transaction:

- If an exception was thrown, the application calls the `javax.transaction.UserTransaction.rollback` method if an exception was thrown during any of the database operations.
- If no exception was thrown, the application calls the `javax.transaction.UserTransaction.commit` method to attempt to commit the transaction after all database operations completed successfully. Calling this method ends the transaction and starts the processing of the operation, causing the WebLogic Enterprise EJB container to call the transaction manager to complete the transaction. The transaction is committed only if all of the participants in the transaction agree to commit.

---

### Listing 1-9 Completing a Transaction

---

```
if(gotException){
    try{
        tx.rollback();
    }catch(Exception e){}
}
elseif{
    tx.commit();
}
```

## Transactions Sample RMI Code

This topic provides a walkthrough of sample code fragments from a class in an RMI application. This topic includes the following sections:

- Importing Packages
- Initializing the UserTransaction Object
- Using JDNI to Return an Object Reference to the UserTransaction Object
- Starting a Transaction

- **Completing a Transaction**

The code fragments demonstrate using the `UserTransaction` object for RMI transactions. For guidelines about using transactions in RMI applications, see Chapter 6, “Transactions in RMI Applications.”

**Note:** These code fragments do not derive from any of the sample applications that ship with WebLogic Enterprise. They merely illustrate the use of the `UserTransaction` object within an RMI application.

## Importing Packages

Listing 1-10 shows importing the necessary packages, including the following packages used to handle transactions:

- `javax.transaction.UserTransaction`. For a list of methods associated with this object, see “UserTransaction Methods” on page 2-24.
- System exceptions. For a list of exceptions, see “Exceptions Thrown by UserTransaction Methods” on page 2-26.

### Listing 1-10 Importing Packages

---

```
import javax.naming.*;
import java.rmi.*;
import javax.transaction.UserTransaction;
import javax.transaction.SystemException;
import javax.transaction.HeuristicMixedException
import javax.transaction.HeuristicRollbackException
import javax.transaction.NotSupportedException
import javax.transaction.RollbackException
import javax.transaction.IllegalStateException
import javax.transaction.SecurityException
```

## Initializing the UserTransaction Object

Listing 1-11 shows initializing an instance of the `UserTransaction` object to null.

### Listing 1-11 Initializing the UserTransaction Object

---

```
UserTransaction tx = null;
```

## Using JNDI to Return an Object Reference to the UserTransaction Object

Listing 1-12 shows searching the JNDI tree to return an object reference to the `UserTransaction` object for the appropriate WebLogic Enterprise domain.

**Note:** Obtaining the object reference begins a conversational state between the application and that object. The conversational state continues until the transaction is completed (committed or rolled back). Once instantiated, RMI objects remain active in memory until they are released (typically during server shutdown). For the duration of the transaction, the WebLogic Enterprise infrastructure does not perform any deactivation or activation.

### Listing 1-12 Performing a JNDI Lookup

---

```
try {  
    Context ctx = getInitialContext();  
    tx = (UserTransaction)ctx.lookup("java:comp/UserTransaction");  
}
```

## Starting a Transaction

Listing 1-13 shows starting a transaction by calling the `javax.transaction.UserTransaction.begin` method. Database operations that occur after this method invocation and prior to completing the transaction exist within the scope of this transaction.

### Listing 1-13 Starting a Transaction

---

```
tx.begin();
```

## Completing a Transaction

Listing 1-14 shows completing the transaction depending on whether an exception was thrown during any of the database operations that were attempted within the scope of this transaction:

- If an exception was thrown, the application calls the `javax.transaction.UserTransaction.rollback` method if an exception was thrown during any of the database operations.
- If no exception was thrown, the application calls the `javax.transaction.UserTransaction.commit` method to attempt to commit the transaction after all database operations completed successfully. Calling this method ends the transaction and starts the processing of the operation, causing WebLogic Enterprise to call the transaction manager to complete the transaction. The transaction is committed only if all of the participants in the transaction agree to commit.

### **Listing 1-14 Completing a Transaction**

---

```
if(gotException){
    try{
        tx.rollback();
    }catch(Exception e){}
}
elseif{
    tx.commit();
}
```

# 2 Transaction Service

This topic includes the following sections:

- About the Transaction Service
- Capabilities and Limitations
- Transaction Service in CORBA Applications
- Transaction Service in EJB Applications
- Transaction Service in RMI Applications
- UserTransaction API

This topic provides the information that programmers need to write transactional applications for the WebLogic Enterprise system. Before you begin, you should read Chapter 1, “Introducing Transactions.”

# About the Transaction Service

WebLogic Enterprise provides a Transaction Service that supports transactions in CORBA, EJB, and RMI applications. The Transaction Service provides:

- An implementation of the CORBAservices Object Transaction Service (OTS) that is described in Chapter 10 of the *CORBAservices: Common Object Services Specification*. This specification defines the interfaces for an object service that provides transactional functions.
- In the WebLogic Enterprise EJB container, an implementation of the transaction services described in the Enterprise JavaBeans Specification 1.1, published by Sun Microsystems, Inc.

For CORBA Java, EJB, and RMI applications, WebLogic Enterprise also provides the `javax.transaction` package, from Sun Microsystems, Inc., which implements the Java Transaction API (JTA) for Java applications. For more information about the JTA, see the Java Transaction API (JTA) Specification (version 1.0.1), published by Sun Microsystems, Inc. For more information about the `UserTransaction` object that applications use to demarcate transaction boundaries, see “UserTransaction API” on page 2-24.

## Capabilities and Limitations

This topic includes the following sections:

- Lightweight Clients with Delegated Commit
- Transaction Propagation (CORBA Only)
- Transaction Integrity
- Transaction Termination
- Flat Transactions
- Interoperability Between Remote Clients and the WebLogic Enterprise Domain

- Intradomain and Interdomain Interoperability
- Network Interoperability
- Relationship of the Transaction Service to Transaction Processing
- Process Failure
- Multithreaded Transaction Client Support
- General Constraints

These sections describe the capabilities and limitations of the Transaction Service that supports CORBA and EJB applications:

## Lightweight Clients with Delegated Commit

A *lightweight client* runs on a single-user, unmanaged desktop system that has irregular availability. Owners may turn their desktop systems off when they are not in use. These single-user, unmanaged desktop systems should not be required to perform network functions such as transaction coordination. In particular, unmanaged systems should not be responsible for ensuring atomicity, consistency, isolation, and durability (ACID) properties across failures for transactions involving server resources. WebLogic Enterprise remote clients are lightweight clients.

The Transaction Service allows lightweight clients to do a delegated commit, which means that the Transaction Service allows lightweight clients to begin and terminate transactions while the responsibility for transaction coordination is delegated to a transaction manager running on a server machine. Client applications do not require a local transaction server. The remote `TransactionCurrent` implementation that CORBA clients use, or the remote implementation of `UserTransaction` that EJB or RMI clients use, delegate the actual responsibility of transaction coordination to transaction manager on the server.

# Transaction Propagation (CORBA Only)

For CORBA applications, the CORBAservices Object Transaction Service specification states that a client can choose to propagate a transaction context either implicitly or explicitly. WebLogic Enterprise *provides* implicit propagation. Explicit propagation *is strongly discouraged*.

**Note:** For EJB and RMI applications, only implicit propagation is supported for clients.

Objects that are related to transaction contexts that are passed around using explicit transaction propagation *should not* be mixed with implicit transaction propagation APIs. It should be noted, however, that explicit propagation does not place any constraints on when transactional methods can be processed. There is no guarantee that all transactional methods will be completed before the transaction is committed.

# Transaction Integrity

Checked transaction behavior provides transaction integrity by guaranteeing that a `commit` will not succeed unless all transactional objects involved in the transaction have completed the processing of their transactional requests. If implicit transaction propagation is used, the Transaction Service *provides* checked transaction behavior that is equivalent to that provided by the request/response interprocess communication models defined by The Open Group. For CORBA applications, for example, the Transaction Service performs `reply` checks, `commit` checks, and `resume` checks, as described in the *CORBAservices Object Transaction Service Specification*.

Unchecked transaction behavior relies completely on the application to provide transaction integrity. If explicit propagation is used, the Transaction Service *does not* provide checked transaction behavior and transaction integrity *is not* guaranteed.

# Transaction Termination

WebLogic Enterprise allows transactions to be terminated *only* by the client that created the transaction.



**Note:** The client may be a server object that requests the services of another object.

## Flat Transactions

WebLogic Enterprise implements the flat transaction model. Nested transactions are *not* supported.

## Interoperability Between Remote Clients and the WebLogic Enterprise Domain

WebLogic Enterprise supports remote clients invoking methods on server objects in *different* WebLogic Enterprise domains in the *same* transaction.

Remote clients with multiple connections to the same WebLogic Enterprise domain *may* make invocations to server objects on these separate connections within the same transaction.

## Intradomain and Interdomain Interoperability

For C++ (but not Java) applications, WebLogic Enterprise supports native clients invoking methods on server objects in the WebLogic Enterprise domain. In addition, WebLogic Enterprise supports server objects invoking methods on other objects in the same or in different processes within the same WebLogic Enterprise domain.

In WebLogic Enterprise applications, transactions can span multiple domains as long as factory-based routing is properly configured across multiple domains. To support transactions across multiple domains, you must configure the `factory_finder.ini` file to identify factory objects that are used in the current (local) domain but that are resident in a different (remote) domain. For more information, see “Configuring Multiple Domains (WebLogic Enterprise System)” in the *Administration Guide*.

# Network Interoperability

A client application can have only one active bootstrap object and `TransactionCurrent` object within a single domain. WebLogic Enterprise does *not* support exporting or importing transactions to or from remote WebLogic Enterprise domains.

However, transactions can encompass multiple domains in a serial fashion. For example, a server with a transaction active in Domain A can communicate with a server in Domain B within the context of that same transaction.

# Relationship of the Transaction Service to Transaction Processing

The Transaction Service relates to various transaction processing servers, interfaces, protocols, and standards in the following ways:

- **Support for BEA Tuxedo ATMI servers.** Servers using the WebLogic Enterprise Transaction Service can make invocations on other BEA Tuxedo Application-to-Transaction Monitor Interface (ATMI) server processes in the same domain. In addition, ATMI services can invoke CORBA objects in both transactional and non-transactional contexts, both within the same domain and across domains via a TDOMAINS gateway. However, WebLogic Enterprise *does not* support remote clients or native clients invoking ATMI services in the WebLogic Enterprise domain.
- **Support for The Open Group XA interface.** The Open Group Resource Managers are resource managers that can be involved in a distributed transaction by allowing their two-phase commit protocol to be controlled via The Open Group XA interface. WebLogic Enterprise supports interaction with The Open Group Resource Managers.
- **Support for the OSI TP protocol.** Open Systems Interconnect Transaction Processing (OSI TP) is the transactional protocol defined by the International Organization for Standardization (ISO). WebLogic Enterprise *does not* support interactions with OSI TP transactions.

- **Support for the LU 6.2 protocol.** Systems Network Architecture (SNA) LU 6.2 is a transactional protocol defined by IBM. WebLogic Enterprise *does not* support interactions with LU 6.2 transactions.
- **Support for the ODMG standard.** ODMG-93 is a standard defined by the Object Database Management Group (ODMG) that describes a portable interface to access Object Database Management Systems. WebLogic Enterprise *does not* support interactions with ODMG transactions.

## Process Failure

The Transaction Service monitors the participants in a transaction for failures and inactivity. The BEA Tuxedo system provides management tools for keeping the application running when failures occur. Because WebLogic Enterprise is built upon the existing BEA Tuxedo transaction management system, it inherits the Tuxedo capabilities for keeping applications running.

## Multithreaded Transaction Client Support

WebLogic Enterprise supports multithreaded clients for non-transactional clients. For transactional clients, WebLogic Enterprise supports only single-threaded client implementation. Clients cannot make transaction requests concurrently in multiple threads.

## General Constraints

The following constraints apply to the Transaction Service:

- In WebLogic Enterprise, a client or a server object *cannot* invoke methods on an object that is infected with (or participating in) another transaction. The method invocation issued by the client or the server will return an exception.
- For CORBA applications, a server application object using transactions from the WebLogic Enterprise Transaction Service library *requires* the TP Framework functionality. For more information about the TP Framework, see “TP Framework” in the *CORBA C++ Programming Reference*.

- For CORBA applications, a return from the `rollback` method on the `Current` object is asynchronous. Similarly, for EJB and RMI applications, a return from the `rollback` method on the `UserTransaction` object is asynchronous.

As a result, the objects that were infected by (or participating in) the rolled back transaction get their states cleared by WebLogic Enterprise *a little later*. Therefore, *no* other client can infect these objects with a different transaction until WebLogic Enterprise clears the states of these objects. This condition exists for a very short amount of time and is typically not noticeable in a production application. A simple workaround for this race condition is to try the appropriate operation after a short (typically a 1-second) delay.

- In WebLogic Enterprise, clients using third-party implementations of the CORBAServices Object Transaction Service (for CORBA applications) or the Java Transaction API (for Java applications) *are not* supported.
- In WebLogic Enterprise CORBA applications, clients may not make one-way method invocations within the context of a transaction to server objects having the `NEVER`, `OPTIONAL`, or `ALWAYS` transaction policies.

No error or exception will be returned to the client because it is a one-way method invocation. However, the method on the server object will not be executed, and an appropriate error message will be written to the log. Clients may make one-way method invocations within the context of a transaction to server objects with the `IGNORE` transaction policy. In this case, the method on the server object will be executed, but not in the context of a transaction. For more information about the transaction policies, see “Server Description File” in the *CORBA Java Programming Reference* or “Implementation Configuration File (ICF)” in the *CORBA C++ Programming Reference*.

# Transaction Service in CORBA Applications

This topic includes the following sections:

- Getting Initial References to the `TransactionCurrent` Object
- CORBA Transaction Service API
- CORBA Transaction Service API Extensions

■ Notes on Using Transactions in WebLogic Enterprise CORBA Applications

These sections describe how WebLogic Enterprise implements the OTS, with particular emphasis on the portion of the CORBAservices Object Transaction Service that is described as implementation-specific. They describe the OTS application programming interface (API) that you use to begin or terminate transactions, suspend or resume transactions, and get information about transactions.

## Getting Initial References to the TransactionCurrent Object

To access the Transaction Service API and the extension to the Transaction Service API as described later in this chapter, an application needs to complete the following operations:

1. Create a `Bootstrap` object. For more information about creating a `Bootstrap` object, see “C++ Bootstrap Object Programming Reference” in the *CORBA C++ Programming Reference*.
2. Invoke the `resolve_initial_reference("TransactionCurrent")` method on the `Bootstrap` object. The invocation returns a standard CORBA object pointer. For a description of this `Bootstrap` object method, see the *CORBA C++ Programming Reference*.
3. If an application requires only the Transaction Service APIs, it should issue an `org.omg.CosTransactions.Current.narrow()` (in Java) or `CosTransactions::Current::_narrow()` (in C++) on the object pointer returned from step 2 above.

If an application requires the Transaction Service APIs with the extensions, it should issue a `com.beasys.Tobj.TransactionCurrent.narrow()` (in Java) or `Tobj::TransactionCurrent::_narrow()` (in C++) on the object pointer returned from step 2 above.

## CORBA Transaction Service API

This topic includes the following sections:

- Data Types
- Exceptions
- Current Interface
- Control Interface
- TransactionalObject Interface

These sections describe the CORBA-based components of the `CosTransactions` modules that WebLogic Enterprise implements to support the Transaction Service. For more information about these components, see Chapter 10 of the *CORBAservices: Common Object Services Specification*.

### Data Types

Listing 2-1 shows the supported data types.

#### Listing 2-1 Data Types Supported by the Transaction Service

---

```
enum Status {  
    StatusActive,  
    StatusMarkedRollback,  
    StatusPrepared,  
    StatusCommitted,  
    StatusRolledBack,  
    StatusUnknown,  
    StatusNoTransaction,  
    StatusPreparing,  
    StatusCommitting,  
    StatusRollingBack  
};  
  
// This information comes from CORBAservices: Common Object  
// Services Specification, p. 10-15. Revised Edition:  
// March 31, 1995. Updated: March 1997. Used with permission by OMG.
```

---

### Exceptions

Listing 2-2 shows the supported exceptions in IDL code.

**Listing 2-2 Exceptions Supported by the Transaction Service**


---

```
// Heuristic exceptions
exception HeuristicMixed {};
exception HeuristicHazard {};

// Other transaction-specific exceptions
exception SubtransactionsUnavailable {};
exception NoTransaction {};
exception InvalidControl {};
exception Unavailable {};
```

---

Table 2-1 describes the exceptions.

**Note:** This information comes from *CORBAservices: Common Object Services Specification*, pages 10-16, 19, 20. Revised Edition: March 31, 1995. Updated: March 1997. Used with permission by OMG.

**Table 2-1 Exceptions Supported by the Transaction Service**

Exception	Description
HeuristicMixed	A request raises this exception to report that a heuristic decision was made and that some relevant updates have been committed and others have been rolled back.
HeuristicHazard	A request raises this exception to report that a heuristic decision was made, that the disposition of all relevant updates is not known, and that for those updates whose disposition is known, either all have been committed or all have been rolled back. Therefore, the <i>HeuristicMixed</i> exception takes priority over the <i>HeuristicHazard</i> exception.
SubtransactionsUnavailable	This exception is raised for the <i>Current</i> interface <i>begin</i> method if the client already has an associated transaction.
NoTransaction	This exception is raised for the <i>Current</i> interface <i>rollback</i> and <i>rollback_only</i> methods if there is no transaction associated with the client thread.

**Table 2-1 Exceptions Supported by the Transaction Service (Continued)**

Exception	Description
InvalidControl	This exception is raised for the Current interface resume method if the parameter is not valid in the current execution environment.
Unavailable	This exception is raised for the Control interface get_terminator and get_coordinator methods if the Control interface cannot provide the requested object.

### Current Interface

The Current interface defines methods that allow a client of the Transaction Service to explicitly manage the association between threads and transactions. The Current interface also defines methods that simplify the use of the Transaction Service for most applications. These methods can be used to begin and end transactions, to suspend and resume transactions, and to obtain information about the current transaction.

The CosTransactions module defines the Current interface (shown in Listing 2-3).

**Listing 2-3 Current Interface idl**

---

```
// Current transaction
interface Current : CORBA::Current {
    void begin()
        raises(SubtransactionsUnavailable);
    void commit(in boolean report_heuristics)
        raises(
            NoTransaction,
            HeuristicMixed,
            HeuristicHazard
        );
    void rollback()
        raises(NoTransaction);
    void rollback_only()
        raises(NoTransaction);
    Status get_status();
    string get_transaction_name();
    void set_timeout(in unsigned long seconds);
    Control get_control();
    Control suspend();
}
```



```
void resume(in Control which)
    raises(InvalidControl);

};

// This information comes from CORBAServices: Common Object
// Services Specification, p. 10-18. Revised Edition:
// March 31, 1995. Updated: November 1997. Used with permission by
// OMG
```

---

Table 2-2 provides a description of the `Current` transaction methods.

**Note:** This information comes from *CORBAServices: Common Object Services Specification*, pages 10-18, 19, 20. Revised Edition: March 31, 1995. Updated: November 1997. Used with permission by OMG.

**Table 2-2 Transaction Methods in the Current Object**

Method	Description
<code>begin</code>	Creates a new transaction. The transaction context of the client thread is modified so that the thread is associated with the new transaction. If the client thread is currently associated with a transaction, the <code>SubtransactionsUnavailable</code> exception is raised. If the client thread cannot be placed in transaction mode due to an error while starting the transaction, the standard system exception <code>INVALID_TRANSACTION</code> is raised. If the call was made in an improper context, the standard system exception <code>BAD_INV_ORDER</code> is raised.

**Table 2-2 Transaction Methods in the Current Object (Continued)**

Method	Description
<code>commit</code>	<p>If there is no transaction associated with the client thread, the <code>NoTransaction</code> exception is raised.</p> <p>If the call was made in an improper context, the standard system exception <code>BAD_INV_ORDER</code> is raised.</p> <p>If the system decides to roll back the transaction, the standard exception <code>TRANSACTION_ROLLEDBACK</code> is raised and the thread's transaction context is set to <code>null</code>.</p> <p>A <code>HeuristicMixed</code> exception is raised to report that a heuristic decision was made and that some relevant updates have been committed and others have been rolled back. A <code>HeuristicHazard</code> exception is raised to report that a heuristic decision was made, and that the disposition of all relevant updates is not known; for those updates whose disposition is known, either all have been committed or all have been rolled back. The <code>HeuristicMixed</code> exception takes priority over the <code>HeuristicHazard</code> exception. If a heuristic exception is raised or the operation completes normally, the thread's transaction exception context is set to <code>null</code>.</p> <p>If the operation completes normally, the thread's transaction context is set to <code>null</code>.</p>
<code>rollback</code>	<p>If there is no transaction associated with the client thread, the <code>NoTransaction</code> exception is raised.</p> <p>If the call was made in an improper context, the standard system exception <code>BAD_INV_ORDER</code> is raised.</p> <p>If the operation completes normally, the thread's transaction context is set to <code>null</code>.</p>
<code>rollback_only</code>	<p>If there is no transaction associated with the client thread, the <code>NoTransaction</code> exception is raised. Otherwise, the transaction associated with the client thread is modified so that the only possible outcome is to roll back the transaction.</p>
<code>get_status</code>	<p>If there is no transaction associated with the client thread, the <code>StatusNoTransaction</code> value is returned. Otherwise, this method returns the status of the transaction associated with the client thread.</p>

**Table 2-2 Transaction Methods in the Current Object (Continued)**

Method	Description
<code>get_transaction_name</code>	If there is no transaction associated with the client thread, an empty string is returned. Otherwise, this method returns a printable string describing the transaction (specifically, the <code>XID</code> as specified by The Open Group). The returned string is intended to support debugging.
<code>set_timeout</code>	<p>This method modifies a state variable associated with the target object that affects the time-out period associated with transactions created by subsequent invocations of the <code>begin</code> method.</p> <p>The initial transaction timeout value is 300 seconds. Calling <code>set_timeout()</code> with an argument value larger than zero specifies a new timeout value. Calling <code>set_timeout()</code> with a zero argument sets the timeout value back to the default of 300 seconds.</p> <p>After calling <code>set_timeout()</code>, transactions created by subsequent invocations of <code>begin</code> are subject to being rolled back if they do not complete before the specified number of seconds after their creation.</p> <p><b>Note:</b> The initial transaction timeout value is 300 seconds. If a transaction is started via <code>AUTOTRAN</code> instead of the <code>begin</code> method, then the timeout value is determined by the <code>TRANTIME</code> value in the WebLogic Enterprise configuration file. For more information, see Chapter 8, “Administering Transactions.”</p>
<code>get_control</code>	If the client is not associated with a transaction, a null object reference is returned. Otherwise, a <code>Control</code> object is returned that represents the transaction context currently associated with the client thread. This object may be given to the <code>resume</code> method to reestablish this context.

**Table 2-2 Transaction Methods in the Current Object (Continued)**

Method	Description
<code>suspend</code>	<p>If the client thread is not associated with a transaction, a null object reference is returned.</p> <p>If the associated transaction is in a state such that the only possible outcome of the transaction is to be rolled back, the standard system exception <code>TRANSACTION_ROLLEDBACK</code> is raised and the client thread becomes associated with no transaction.</p> <p>If the call was made in an improper context, the standard system exception <code>BAD_INV_ORDER</code> is raised. The caller's state with respect to the transaction is not changed.</p> <p>Otherwise, an object is returned that represents the transaction context currently associated with the client thread. The same client can subsequently give this object to the <code>resume</code> method to reestablish this context. In addition, the client thread becomes associated with no transaction.</p> <p><b>Note:</b> As defined in The Common Object Request Broker: Architecture and Specification, Revision 2.2, February 1998, the standard system exception <code>TRANSACTION_ROLLEDBACK</code> indicates that the transaction associated with the request has already been rolled back or has been marked to roll back. Thus, the requested method either could not be performed or was not performed because further computation on behalf of the transaction would be fruitless.</p>

**Table 2-2 Transaction Methods in the Current Object (Continued)**

Method	Description
<code>resume</code>	<p>If the client thread is already associated with a transaction which is in a state such that the only possible outcome of the transaction is to be rolled back, the standard system exception <code>TRANSACTION_ROLLEDBACK</code> is raised and the client thread becomes associated with no transaction.</p> <p>If the call was made in an improper context, the standard system exception <code>BAD_INV_ORDER</code> is raised.</p> <p>If the system is unable to resume the global transaction because the caller is currently participating in work outside any global transaction with one or more resource managers, the standard system exception <code>INVALID_TRANSACTION</code> is raised.</p> <p>If the parameter is a <code>null</code> object reference, the client thread becomes associated with no transaction. If the parameter is valid in the current execution environment, the client thread becomes associated with that transaction (in place of any previous transaction). Otherwise, the <code>InvalidControl</code> exception is raised.</p> <p><b>Note:</b> See <code>suspend</code> for a definition of the standard system exception <code>TRANSACTION_ROLLEDBACK</code>.</p>

## Control Interface

The `Control` interface allows a program to explicitly manage or propagate a transaction context. An object that supports the `Control` interface is implicitly associated with one specific transaction.

Listing 2-4 shows the `Control` interface, which is defined in the `CosTransactions` module.

### Listing 2-4 Control Interface

```
interface Control {
    Terminator get_terminator()
        raises(Unavailable);
    Coordinator get_coordinator()
```

```
        raises(Unavailable);
    };

    // This information comes from CORBAservices: Common Object
    // Services Specification, p. 10-21. Revised Edition:
    // March 31, 1995. Updated: November 1997. Used with permission by
    // OMG.
```

---

The Control interface is used only in conjunction with the suspend and resume methods.

### TransactionalObject Interface

The `org.omg.CosTransactions.TransactionalObject` interface (in Java) or `CosTransactions::TransactionalObject` (in C++) is used by an object to indicate that it is transactional. By supporting this interface, an object indicates that it wants the transaction context associated with the client thread to be propagated on requests to the object. *However, this interface is no longer needed.* For details on transaction policies that need to be set to infect objects with transactions, see “Server Description File” in the *CORBA Java Programming Reference* or “Implementation Configuration File (ICF)” in the *CORBA C++ Programming Reference*.

The `CosTransactions` module defines the `TransactionalObject` interface (shown in Listing 2-5). The `org.omg.CosTransactions.TransactionalObject` interface defines no methods. It is simply a marker.

#### Listing 2-5 TransactionalObject Interface

---

```
interface TransactionalObject {
};

    // This information comes from CORBAservices: Common Object
    // Services Specification, p. 10-30. Revised Edition:
    // March 31, 1995. Updated: November 1997. Used with permission by
    // OMG.
```

---

## Other CORBAServices Object Transaction Service Interfaces

All other CORBAServices Object Transaction Service interfaces are *not* supported. Note that the `Current` interface described earlier is supported only if it has been obtained from the `Bootstrap` object. The `Control` interface described earlier is supported only if it has been obtained using the `get_control` and the `suspend` methods on the `Current` object.

## CORBA Transaction Service API Extensions

This topic describes specific extensions to the CORBAServices Transaction Service API described earlier. The APIs in this topic enable an application to open or close an Open Group Resource Manager.

The following APIs help facilitate participation of resource managers in a distributed transaction by allowing their two-phase commit protocol to be controlled via The Open Group XA interface.

The following definitions and interfaces are defined in the `com.beasys.Tobj` module (in Java) or `Tobj` module (in C++).

### Exception

The following exception is supported:

```
exception RMfailed {};
```

A request raises this exception to report that an attempt to open or close a resource manager failed.

### TransactionCurrent Interface

This interface supports all the methods of the `Current` interface in the `CosTransactions` module and is described in “Java Bootstrap Object Programming Reference” in the CORBA Java Programming Reference or in “C++ Bootstrap Object Programming Reference” in the CORBA C++ Programming Reference. Additionally, this interface supports APIs to open and close the resource manager.

Listing 2-6 shows the `TransactionCurrent` interface, which is defined in the `Tobj` module.

### Listing 2-6 TransactionCurrent Interface

---

```
Interface TransactionCurrent: CosTransactions::Current {  
    void open_xa_rm()  
        raises(RMfailed);  
    void close_xa_rm()  
        raises(Rmfailed);  
}
```

---

Table 2-3 describes APIs that are specific to the resource manager. For more information about these APIs, see the *CORBA Java Programming Reference* or the *CORBA C++ Programming Reference*.

**Table 2-3 Resource Manager APIs for the Current Interface**

Method	Description
open_xa_rm	<p>This method opens The Open Group Resource Manager to which this process is linked. A <code>RMfailed</code> exception is raised if there is a failure while opening the resource manager.</p> <p>Any attempts to invoke this method by remote clients or the native clients raises the standard system exception <code>NO_IMPLEMENT</code>.</p>
close_xa_rm	<p>This method closes The Open Group Resource Manager to which this process is linked. An <code>RMfailed</code> exception is raised if there is a failure while closing the resource manager. A <code>BAD_INV_ORDER</code> standard system exception is raised if the function was called in an improper context (for example, the caller is in transaction mode).</p> <p>Any attempts by the remote clients or the native clients to invoke this method raises the standard system exception <code>NO_IMPLEMENT</code>.</p>

## Notes on Using Transactions in WebLogic Enterprise CORBA Applications

Consider the following guidelines when integrating transactions into your WebLogic Enterprise CORBA client/server applications:

- Nested transactions are not permitted in the WebLogic Enterprise system.



You cannot start a new transaction if an existing transaction is already active. (You may start a new transaction if you first suspend the existing one; however, the object that suspends the transaction is the only object that can subsequently resume the transaction.)

- The object that starts a transaction is the only entity that can end the transaction. (In a strict sense, the object can be the client application, the TP Framework, or an object managed by the server application.) An object that is invoked within the scope of a transaction may suspend and resume the transaction (and while the transaction is suspended, the object can start and end other transactions). However, you cannot end a transaction in an object unless you began the transaction there.
- WebLogic Enterprise does not support concurrent transactions. Objects can be involved with only one transaction at one time. An object is involved in a transaction for the duration of the entire transaction, and is available to be involved in a different transaction only after the current transaction is completed.
- WebLogic Enterprise does not queue requests to objects that are currently involved in a transaction. If a non-transactional client application attempts to invoke an operation on an object that is currently in a transaction, the client application receives the following error message:

### Java

```
org.omg.CORBA.OBJ_ADAPTER
```

### C++

```
CORBA::OBJ_ADAPTER
```

If a client that is in a transaction attempts to invoke an operation on an object that is currently in a different transaction, the client application receives the following error message:

### Java

```
org.omg.CORBA.INVALID_TRANSACTION
```

### C++

```
CORBA::INVALID_TRANSACTION
```

- For transaction-bound objects, consider doing all state handling in the `com.beasys.Tobj_Servant.deactivate_object` method (in Java) or `Tobj_ServantBase::deactivate_object()` operation (in C++). This makes

it easier for the object to handle its state properly, because the outcome of the transaction is known at the time that `deactivate_object()` is invoked.

- For method-bound objects that have several operations, but only a few that affect the object's durable state, consider doing the following:
  - Assign the `optional` transaction policy.
  - Scope any write operations within a transaction, by making invocations on the `TransactionCurrent` object.

If the object is invoked outside a transaction, the object does not incur the overhead of scoping a transaction for reading data. This way, regardless of whether the object is invoked within a transaction, all the object's write operations are handled transactionally.

- Transaction rollbacks are asynchronous. Therefore, it is possible for an object to be invoked while its transactional context is still active. If you try to invoke such an object, you receive an exception.
- If an object with the `always` transaction policy is involved in a transaction that is started by the WebLogic Enterprise system, and not the client application, note the following:
  - If the server application marks the transaction for rollback only and the server throws a CORBA exception, the client application receives the CORBA exception.
  - If the server application marks the transaction for rollback only and the server does *not* throw a CORBA exception, the client application receives the `OBJ_ADAPTER` exception. In this case, the WebLogic Enterprise system automatically rolls back the transaction. However, the client application is completely unaware that a transaction has been scoped in the WebLogic Enterprise domain.
- If the client application initiates a transaction, and the server application marks the transaction for a rollback, one of the following occurs:
  - If the server throws a CORBA exception, the client application receives a CORBA exception.
  - If the server does *not* throw a CORBA exception, the client application receives the `TRANSACTION_ROLLEDBACK` exception.

# Transaction Service in EJB Applications

The WebLogic Enterprise EJB container provides a Transaction Service that supports the two types of transactions in WebLogic Enterprise EJB applications:

- **Container-managed transactions.** In container-managed transactions, the WebLogic Enterprise EJB container manages the transaction demarcation. Transaction attributes in the EJB deployment descriptor determine how the WebLogic Enterprise EJB container handles transactions with each method invocation.
- **Bean-managed transactions.** In bean-managed transactions, the EJB manages the transaction demarcation. The EJB makes explicit method invocations on the `UserTransaction` object to begin, commit, and roll back transactions. For more information about `UserTransaction` methods, see “UserTransaction API” on page 2-24.

For an introduction to transaction management in EJB applications, see “Transactions in WebLogic Enterprise EJB Applications” on page 1-8, and “Transactions Sample EJB Code” on page 1-25.

# Transaction Service in RMI Applications

WebLogic Enterprise provides a Transaction Service that supports transactions in WebLogic Enterprise RMI applications. In RMI applications, the client or server application makes explicit method invocations on the `UserTransaction` object to begin, commit, and roll back transactions.

For more information about `UserTransaction` methods, see “UserTransaction API” on page 2-24. For an introduction to transaction management in RMI applications, see “Transactions in WebLogic Enterprise RMI Applications” on page 1-11, and “Transactions Sample RMI Code” on page 1-27.

# UserTransaction API

This topic includes the following sections:

- UserTransaction Methods
- Exceptions Thrown by UserTransaction Methods

WebLogic Enterprise provides the `javax.transaction` package, from Sun Microsystems, Inc., which implements the Java Transaction API (JTA) for Java applications. The `javax.UserTransaction` interface supports transaction management for CORBA Java applications as well as for bean-managed transactions in EJB applications. For more information about the JTA, see the Java Transaction API (JTA) Specification (version 1.0.1) published by Sun Microsystems, Inc. For a detailed description of the `javax.transaction` interface, see the package description in the *WebLogic Enterprise Javadoc*.

## UserTransaction Methods

Table 2-4 describes the methods in the `UserTransaction` object.

**Table 2-4 Methods in the UserTransaction Object**

Method Name	Description
<code>begin</code>	Starts a transaction on the current thread.
<code>commit</code>	Commits the transaction associated with the current thread.

**Table 2-4 Methods in the UserTransaction Object (Continued)**

Method Name	Description
<code>getStatus</code>	Returns the transaction status, or <code>STATUS_NO_TRANSACTION</code> if no transaction is associated with the current thread.  One of the following values: <ul style="list-style-type: none"><li>■ <code>STATUS_ACTIVE</code></li><li>■ <code>STATUS_COMMITTED</code></li><li>■ <code>STATUS_COMMITTING</code></li><li>■ <code>STATUS_MARKED_ROLLBACK</code></li><li>■ <code>STATUS_NO_TRANSACTION</code></li><li>■ <code>STATUS_PREPARED</code></li><li>■ <code>STATUS_PREPARING</code></li><li>■ <code>STATUS_ROLLEDBACK</code></li><li>■ <code>STATUS_ROLLING_BACK</code></li><li>■ <code>STATUS_UNKNOWN</code></li></ul>
<code>rollback</code>	Rolls back the transaction associated with the current thread.
<code>setRollbackOnly</code>	Marks the transaction associated with the current thread so that the only possible outcome of the transaction is to roll it back.
<code>setTransactionTimeout</code>	Specifies the timeout value for the transactions started by the current thread with the <code>begin</code> method. If an application has not called the <code>begin</code> method, then the Transaction Service uses a default value for the transaction timeout.

# Exceptions Thrown by UserTransaction Methods

Table 2-5 describes exceptions thrown by methods of the `UserTransaction` object.

**Table 2-5 Exceptions Thrown by UserTransaction Methods**

Exception	Description
HeuristicMixedException	Thrown to indicate that a heuristic decision was made and that some relevant updates have been committed while others have been rolled back.
HeuristicRollbackException	Thrown to indicate that a heuristic decision was made and that some relevant updates have been rolled back.
NotSupportedException	Thrown when the requested operation is not supported (such as a nested transaction).
RollbackException	Thrown when the transaction has been marked for rollback only or the transaction has been rolled back instead of committed.
IllegalStateException	Thrown if the current thread is not associated with a transaction.
SecurityException	Thrown to indicate that the thread is not allowed to commit the transaction.
SystemException	Thrown by the transaction manager to indicate that it has encountered an unexpected error condition that prevents future transaction services from proceeding.

# 3 Transactions in CORBA Server Applications

This topic includes the following sections:

- Integrating Transactions in a WebLogic Enterprise Client and Server Application
- Transactions and Object State Management
- User-defined Exceptions

These sections describe how to integrate transactions into a WebLogic Enterprise server application. Before you begin, you should read Chapter 1, “Introducing Transactions.”

# Integrating Transactions in a WebLogic Enterprise Client and Server Application

This topic includes the following sections:

- Transaction Support in CORBA Applications
- Making an Object Automatically Transactional
- Enabling an Object to Participate in a Transaction
- Preventing an Object from Being Invoked While a Transaction Is Scoped
- Excluding an Object from an Ongoing Transaction
- Assigning Policies
- Using an XA Resource Manager
- Opening an XA Resource Manager
- Closing an XA Resource Manager

## Transaction Support in CORBA Applications

WebLogic Enterprise supports transactions in the following ways:

- The client or the server application can begin and end transactions explicitly by using calls on the `TransactionCurrent` object. For details about the `TransactionCurrent` object, see Chapter 4, “Transactions in CORBA Client Applications.”
- You can assign transactional policies to an object’s interface so that when the object is invoked, the WebLogic Enterprise system can start a transaction automatically for that object, if a transaction has not already been started, and commit or roll back the transaction when the method invocation is complete. You use transactional policies on objects in conjunction with an XA Resource Manager and database when you want to delegate all the transaction commit and rollback responsibilities to that resource manager.



- Objects involved in a transaction can force a transaction to be rolled back. That is, after an object has been invoked within the scope of a transaction, the object can invoke `rollback_only()` on the `TransactionCurrent` object to mark the transaction for rollback only. This prevents the current transaction from being committed. An object may need to mark a transaction for rollback if an entity, typically a database, is otherwise at risk of being updated with corrupt or inaccurate data.
- Objects involved in a transaction can be kept in memory from the time they are first invoked until the moment when the transaction is ready to be committed or rolled back. In the case of a transaction that is about to be committed, these objects are polled by the WebLogic Enterprise system immediately before the resource managers prepare to commit the transaction. In this sense, polling means invoking the object's `com.beasys.Tobj_Servant.deactivate_object` method (in Java) or `Tobj_ServantBase::deactivate_object()` operation (in C++) and passing a reason value.

When an object is polled, the object may veto the current transaction by invoking `rollback_only()` on the `TransactionCurrent` object. In addition, if the current transaction is to be rolled back, objects have an opportunity to skip any writes to a database. If no object vetoes the current transaction, the transaction is committed.

The following sections explain how you can use object activation policies and transaction policies to determine the transactional behavior you want in your objects. Note that these policies apply to an interface and, therefore, to all operations on all objects implementing that interface.

**Note:** If a server application manages an object that you want to be able to participate in a transaction, the `Server` object for that application must invoke the `com.beasys.Tobj.TP.open_xa_rm` and `com.beasys.Tobj.TP.close_xa_rm` methods (in Java), or the `TP::open_xa_rm()` and `TP::close_xa_rm()` operations (in C++). For more information about database connections, see “Opening an XA Resource Manager” on page 3-8.

## Making an Object Automatically Transactional

The WebLogic Enterprise system provides the `always` transactional policy, which you can define on an object's interface to have the WebLogic Enterprise system start a transaction automatically when that object is invoked and a transaction has not already been scoped. When an invocation on that object is completed, the WebLogic Enterprise system commits or rolls back the transaction automatically. Neither the server application, nor the object implementation, needs to invoke the `TransactionCurrent` object in this situation; the WebLogic Enterprise system automatically invokes the `TransactionCurrent` object on behalf of the server application.

Assign the `always` transactional policy to an object's interface when:

- The object writes to a database and you want all the database commit or rollback responsibilities delegated to an XA Resource Manager whenever this object is invoked.
- You want to give the client application the opportunity to include the object in a larger transaction that encompasses invocations on multiple objects, and the invocations must all succeed or be rolled back if any one invocation fails.

If you want an object to be automatically transactional, assign the following policies to that object's interface in the XML-based Server Description File (in Java) or Implementation Configuration File (in C++):

Activation Policies	Transaction Policy
<ul style="list-style-type: none"><li>■ <code>process</code></li><li>■ <code>method</code></li><li>■ <code>transaction</code></li></ul>	<code>always</code>

**Note:** Database cursors cannot span transactions. However, in C++, the `CourseSynopsisEnumerator` object in the WebLogic Enterprise University sample applications uses a database cursor to find matching course synopses from the University database. Because database cursors cannot span transactions, the `activate_object()` operation on the `CourseSynopsisEnumerator` object reads all matching course synopses into memory. Note that the cursor is managed by an iterator class and is thus not visible to the `CourseSynopsisEnumerator` object.

## Enabling an Object to Participate in a Transaction

If you want an object to be able to be invoked within the scope of a transaction, you can assign the `optional` transaction policies to that object's interface. The `optional` transaction policy may be appropriate for an object that does not perform any database write operations, but that you want to have the ability to be invoked during a transaction.

You can use the following policies, when they are specified in the XML-based Server Description File (in Java) or Implementation Configuration File (in C++) for that object's interface, to make an object optionally transactional:

Activation Policies	Transaction Policy
<ul style="list-style-type: none"><li>■ <code>process</code></li><li>■ <code>method</code></li><li>■ <code>transaction</code></li></ul>	<code>optional</code>

When the transaction policy is `optional`, if the `AUTOTRAN` parameter is enabled in the application's `UBBCONFIG` file, the implementation is transactional. Servers containing transactional objects must be configured within a group associated with an XA-compliant Resource Manager.

If the object does perform database write operations, and you want the object to be able to participate in a transaction, assigning the `always` transactional policy is generally a better choice. However, if you prefer, you can use the `optional` policy and encapsulate any write operations within invocations on the `TransactionCurrent` object. That is, within your operations that write data, scope a transaction around the write statements by invoking the `TransactionCurrent` object to, respectively, begin and commit or roll back the transaction, if the object is not already scoped within a transaction. This ensures that any database write operations are handled transactionally. This also introduces a performance efficiency: if the object is not invoked within the scope of a transaction, all the database read operations are nontransactional, and, therefore, more streamlined.

**Note:** When choosing the transaction policies to assign to your objects, make sure you are familiar with the requirements of the XA Resource Manager you are using. For example, some XA Resource Managers (such as the Oracle 7 Transaction Manager Server) require that any object participating in a

transaction scope their database read operations, in addition to write operations, within a transaction (you can still scope your own transactions, however). Other resource managers, such as Oracle8i, do not require a transaction context for read and write operations. If an application attempts a write operation without a transaction context, Oracle8i will start a local transaction implicitly, in which case the application needs to commit the local transaction explicitly.

## Preventing an Object from Being Invoked While a Transaction Is Scoped

In many cases, it may be critical to exclude an object from a transaction. If such an object is invoked during a transaction, the object returns an exception, which may cause the transaction to be rolled back. The WebLogic Enterprise system provides the `never` transaction policy, which you can assign to an object's interface to specifically prevent that object from being invoked within the course of a transaction, even if the current transaction is suspended.

This transaction policy is appropriate for objects that write durable state to disk that cannot be rolled back, such as for an object that writes data to a disk that is not managed by an XA Resource Manager. Having this capability in your client/server application is crucial if the client application does not or cannot know if some of its invocations are causing a transaction to be scoped. Therefore, if a transaction is scoped, and an object with this policy is invoked, the transaction can be rolled back.

To prevent an object from being invoked while a transaction is scoped, assign the following policies to that object's interface in the XML-based Server Description File (in Java) or Implementation Configuration File (in C++):

Activation Policies	Transaction Policy
■ <code>process</code>	<code>never</code>
■ <code>method</code>	

## Excluding an Object from an Ongoing Transaction

In some cases, it may be appropriate to permit an object to be invoked during the course of a transaction but also keep that object from being a part of the transaction. If such an object is invoked during a transaction, the transaction is automatically suspended. After the invocation on the object is completed, the transaction is automatically resumed. The WebLogic Enterprise system provides the `ignore` transaction policy for this purpose.

The `ignore` transaction policy may be appropriate for an object such as a factory that typically does not write data to disk. By excluding the factory from the transaction, the factory can be available to other client invocations during the course of a transaction. In addition, using this policy can introduce an efficiency into your server application because it minimizes the overhead of invoking objects transactionally.

To prevent any transaction from being propagated to an object, assign the following policies to that object's interface in the Server Description File (in Java) or Implementation Configuration File (in C++):

Activation Policies	Transaction Policy
<ul style="list-style-type: none"><li>■ <code>process</code></li><li>■ <code>method</code></li></ul>	<code>ignore</code>

## Assigning Policies

For information about how to create a Server Description File (in Java) or Implementation Configuration File (in C++) and specify policies on objects, see “Step 5: Define the object activation and transaction policies” in “Steps for Creating a WebLogic Enterprise Server Application” in *Creating Java Server Applications*, or “Step 4: Define the in-memory behavior of objects” in “Steps for Creating a WebLogic Enterprise Server Application” in *Creating C++ Server Applications*.

## Using an XA Resource Manager

The Transaction Manager Server (TMS) handles object state data automatically. For example, the XA Bankapp sample C++ application in the `drive:\M3dir\samples\corba\bankapp_java\XA` directory uses the Oracle TMS as an example of a relational database management service (RDBMS).

Using any XA Resource Manager imposes specific requirements on how different objects managed by the server application may read and write data to that database, including the following:

- Some XA Resource Managers, such as Oracle7, require that all database operations be scoped within a transaction. This means that all method invocations on the `DBAccess` object need to be scoped within a transaction because this object reads from a database. The transaction can be started either by the client or by the WebLogic Enterprise system.

Other XA Resource Managers, such as Oracle8i, do not require a transaction context for read and write operations. If an application attempts a write operation without a transaction context, Oracle8i will start a local transaction implicitly, in which case the application needs to commit the local transaction explicitly.

- When a transaction is committed or rolled back, the XA Resource Manager automatically handles the durable state implied by the commit or rollback. That is, if the transaction is committed, the XA Resource Manager ensures that all database updates are made permanent. Likewise, if there is a rollback, the XA Resource Manager automatically restores the database to its state prior to the beginning of the transaction.

This characteristic of XA Resource Managers actually makes the design problems associated with handling object state data in the event of a rollback much simpler. Transactional objects can always delegate the commit and rollback responsibilities to the XA Resource Manager, which greatly simplifies the task of implementing a server application.

## Opening an XA Resource Manager

This section describes how to open the XA Resource Manager in Java and C++.

## Opening an XA Resource Manager in Java

If an object's interface has the `always` or `optional` transaction policy, you must invoke the `com.beasys.Tobj.TP.open_xa_rm` method in the `com.beasys.Tobj.Server.initialize` method in the `Server` object that supports this object. You must build a special version of the `JavaServer` by using the `buildXAJS` command, if your object performs database operations.

In the `SERVERS` section of the application's `UBBCONFIG` file, you must use the `JavaServerXA` element in place of `JavaServer` to associate the XA Resource Manager with a specified server group. (`JavaServer` uses the null RM.)

The resource manager is opened using the information provided in the `OPENINFO` parameter, which is in the `GROUPS` section of the `UBBCONFIG` file. Note that the default version of the `com.beasys.Tobj.Server.initialize` method automatically opens the resource manager.

If you have an object that participates in a transaction but does not actually perform database operations (the object typically has the `optional` transaction policy), you still need to include an invocation to the `com.beasys.Tobj.TP.open_xa_rm` method.

## Opening an XA Resource Manager in C++

If an object's interface has the `always` or `optional` transaction policy, you must invoke the `TP::open_xa_rm()` operation in the `Server::initialize()` operation in the `Server` object. The resource manager is opened using the information provided in the `OPENINFO` parameter, which is in the `GROUPS` section of the `UBBCONFIG` file. Note that the default version of the `Server::initialize()` operation automatically opens the resource manager.

If you have an object that does not write data to disk and that participates in a transaction—the object typically has the `optional` transaction policy—you still need to include an invocation to the `TP::open_xa_rm()` operation. In that invocation, specify the `NULL` resource manager.

## Closing an XA Resource Manager

If your Server object's `com.beasys.Tobj.Server.initialize` method (in Java) or `Server::initialize()` operation (in C++) opens an XA Resource Manager, you must include the following invocation in the `com.beasys.Tobj.Server.release` method (in Java) or `Server::release()` operation (in C++):

### Java

```
com.beasys.Tobj.TP.close_xa_rm();
```

### C++

```
TP::close_xa_rm();
```

## Transactions and Object State Management

This topic includes the following sections:

- Delegating Object State Management to an XA Resource Manager
- Waiting Until Transaction Work Is Complete Before Writing to the Database

If you need transactions in your WebLogic Enterprise client and server application, you can integrate transactions with object state management in a few different ways. In general, the WebLogic Enterprise system can automatically scope the transaction for the duration of an operation invocation without requiring you to make any changes to your application's logic or the way in which the object writes durable state to disk.

## Delegating Object State Management to an XA Resource Manager

Using an XA Resource Manager, such as Oracle, generally simplifies the design problems associated with handling object state data in the event of a rollback. (The Oracle Resource Manager is used in the WebLogic Enterprise University sample C++ applications). Transactional objects can always delegate the commit and rollback



responsibilities to the XA Resource Manager, which greatly simplifies the task of implementing a server application. This means that process- or method-bound objects involved in a transaction can write to a database during transactions, and can depend on the resource manager to undo any data written to the database in the event of a transaction rollback.

## Waiting Until Transaction Work Is Complete Before Writing to the Database

The `transaction` activation policy is a good choice for objects that maintain state in memory that you do not want written, or that cannot be written, to disk until the transaction work is complete. When you assign the `transaction` activation policy to an object, the object:

- Is brought into memory when it is first invoked within the scope of a transaction.
- Remains in memory until the transaction is either committed or rolled back.

When the transaction work is complete, the WebLogic Enterprise system invokes each transaction-bound object's `com.beasys.Tobj_Servant.deactivate_object` method (in Java) or `Tobj_ServantBase::deactivate_object()` operation (in C++), passing a reason code that can be either `DR_TRANS_COMMITTING` or `DR_TRANS_ABORTED`. If the variable is `DR_TRANS_COMMITTING`, the object can invoke its database write operations. If the variable is `DR_TRANS_ABORTED`, the object skips its write operations.

## When to Assign the Transaction Activation Policy

Assigning the `transaction` activation policy to an object may be appropriate in the following situations:

- You want the object to write its persistent state to disk at the time that the transaction work is complete.

This introduces a performance efficiency because it reduces the number of database write operations that may need to be rolled back.

- You want to provide the object with the ability to veto a transaction that is about to be committed.

If the WebLogic Enterprise system passes the reason `DR_TRANS_COMMITTING`, the object can, if necessary, invoke `rollback_only()` on the `TransactionCurrent` object. Note that if you do make an invocation to `rollback_only()` from within the `com.beasys.Tobj_Servant.deactivate_object` method (in Java) or `Tobj_ServantBase::deactivate_object()` operation (in C++), then `deactivate_object()` is not invoked again.

- You want to provide the object with the ability to perform batch updates.
- You have an object that is likely to be invoked multiple times during the course of a single transaction, and you want to avoid the overhead of continually activating and deactivating the object during that transaction.

## Transaction Policies to Use with the Transaction Activation Policy

To give an object the ability to wait until the transaction is committing before writing to a database, assign the following policies to that object's interface in the XML-based Server Description File (in Java) or Implementation Configuration File (in C++):

Activation Policy	Transaction Policy
<code>transaction</code>	<code>always or optional</code>

**Note:** Transaction-bound objects cannot start a transaction or invoke other objects from inside the `com.beasys.Tobj_Servant.deactivate_object` method (in Java) or `Tobj_ServantBase::deactivate_object()` operation (in C++). The only valid invocations transaction-bound objects can make inside `deactivate_object()` are write operations to the database.

Also, if you have an object that is involved in a transaction, the Server object that manages that object must include invocations to open and close the XA Resource Manager, even if the object does not write any data to disk. (If you have a transactional object that does not write data to disk, you specify the `NULL` resource manager.) For more information about opening and closing an XA Resource Manager, see “Opening an XA Resource Manager” on page 3-8 and “Closing an XA Resource Manager” on page 3-10.

# User-defined Exceptions

This topic includes the following sections:

- About User-defined Exceptions
- Defining the Exception
- Throwing the Exception

## About User-defined Exceptions

Including a user-defined exception in a WebLogic Enterprise client/server application involves the following steps:

1. In your OMG IDL file, define the exception and specify the operations that can use it.
2. In the implementation file, include code that throws the exception.
3. In the client application source file, include code that catches and handles the exception.

For example, the Transactions sample C++ application includes an instance of a user-defined exception, `TooManyCredits`. This exception is thrown by the server application when the client application tries to register a student for a course, and the student has exceeded the maximum number of courses for which he or she can register. When the client application catches this exception, the client application rolls back the transaction that registers a student for a course. This section explains how you can define and implement user-defined exceptions in your WebLogic Enterprise client/server application, using the `TooManyCredits` exception as an example.

## Defining the Exception

In the OMG IDL file for your client/server application:

1. Define the exception and define the data sent with the exception. For example, the `TooManyCredits` exception is defined to pass a short integer representing the maximum number of credits for which a student can register. Therefore, the definition for the `TooManyCredits` exception contains the following OMG IDL statements:

```
exception TooManyCredits
{
    unsigned short maximum_credits;
};
```

2. In the definition of the operations that throw the exception, include the exception. The following example shows the OMG IDL statements for the `register_for_courses()` operation on the Registrar interface:

```
NotRegisteredList register_for_courses(
    in StudentId      student,
    in CourseNumberList courses
) raises (
    TooManyCredits
);
```

## Throwing the Exception

In the implementation of the operation that uses the exception, write the code that throws the exception, as in the following C++ example.

```
if ( ... ) {
    UniversityZ::TooManyCredits e;
    e.maximum_credits = 18;
    throw e;
```

## How the Transactions University Sample Application Works (C++ Only)

This topic includes the following sections:

- About the Transactions University Sample Application

- Transactional Model Used by the Transactions University Sample Application
- Object State Considerations for the University Server Application
- Configuration Requirements for the Transactions Sample Application

## About the Transactions University Sample Application

To implement the student registration process, the Transactions sample application does the following:

- The client application obtains a reference to the `TransactionCurrent` object from the `Bootstrap` object.
- When the student submits the list of courses for which he or she wants to register, the client application:
  - a. Begins a transaction by invoking the `Current::begin()` operation on the `TransactionCurrent` object.
  - b. Invokes the `register_for_courses()` operation on the `Registrar` object, passing a list of courses.
- The `register_for_courses()` operation on the `Registrar` object processes the registration request by executing a loop that does the following iteratively for each course in the list:
  - a. Checks to see how many credits the student is already registered for.
  - b. Adds the course to the list of courses for which the student is registered.

The `Registrar` object checks for the following potential problems, which prevent the transaction from being committed:

- The student is already registered for the course.
  - A course in the list does not exist.
  - The student exceeds the maximum credits allowed.
- As defined in the application's OMG IDL, the `register_for_courses()` operation returns a parameter to the client application, `NotRegisteredList`, which contains a list of the courses for which the registration failed.

If the `NotRegisteredList` value is empty, the client application commits the transaction.

If the `NotRegisteredList` value contains any courses, the client application queries the student to indicate whether he or she wants to complete the registration process for the courses for which the registration succeeded. If the user chooses to complete the registration, the client application commits the transaction. If the user chooses to cancel the registration, the client application rolls back the transaction.

- If the registration for a course has failed because the student exceeds the maximum number of credits he or she can take, the `Registrar` object returns a `TooManyCredits` exception to the client application, and the client application rolls back the entire transaction.

## Transactional Model Used by the Transactions University Sample Application

The basic design rationale for the Transactions sample application is to handle course registrations in groups, as opposed to one at a time. This design helps to minimize the number of remote invocations on the `Registrar` object.

In implementing this design, the Transactions sample application shows one model of the use of transactions, which were described in “Integrating Transactions in a WebLogic Enterprise Client and Server Application” on page 3-2. The model is as follows:

- The client begins a transaction by invoking the `begin()` operation on the `TransactionCurrent` object, followed by making an invocation to the `register_for_courses()` operation on the `Registrar` object.

The `Registrar` object registers the student for the courses for which it can, and then returns a list of courses for which the registration process was unsuccessful. The client application can choose to commit the transaction or roll it back. The transaction encapsulates this conversation between the client and the server application.

- The `register_for_courses()` operation performs multiple checks of the University database. If any one of those checks fail, the transaction can be rolled back.

## Object State Considerations for the University Server Application

Because the Transactions University sample application is transactional, the University server application generally needs to consider the implications on object state, particularly in the event of a rollback. In cases where there is a rollback, the server application must ensure that all affected objects have their durable state restored to the proper state.

Because the `Registrar` object is being used for database transactions, a good design choice for this object is to make it transactional (assign the `always` transaction policy to this object's interface). If a transaction has not already been scoped when this object is invoked, the WebLogic Enterprise system will start a transaction automatically.

By making the `Registrar` object automatically transactional, all database write operations performed by this object will always be done within the scope of a transaction, regardless of whether the client application starts one. Since the server application uses an XA Resource Manager, and since the object is guaranteed to be in a transaction when the object writes to a database, the object does not have any rollback or commit responsibilities because the XA Resource Manager takes responsibility for these database operations on behalf of the object.

The `RegistrarFactory` object, however, can be excluded from transactions because this object does not manage data that is used during the course of a transaction. By excluding this object from transactions, you minimize the processing overhead implied by transactions.

### Object Policies Defined for the Registrar Object

To make the `Registrar` object transactional, the ICF file specifies the `always` transaction policy for the `Registrar` interface. Therefore, in the Transaction sample application, the ICF file specifies the following object policies for the `Registrar` interface:

Activation Policy	Transaction Policy
<code>process</code>	<code>always</code>

## Object Policies Defined for the RegistrarFactory Object

To exclude the `RegistrarFactory` object from transactions, the ICF file specifies the `ignore` transaction policy for the `Registrar` interface. Therefore, in the Transaction sample application, the ICF file specifies the following object policies for the `RegistrarFactory` interface:

Activation Policy	Transaction Policy
<code>process</code>	<code>ignore</code>

## Using an XA Resource Manager in the Transactions Sample Application

The Transactions sample application uses the Oracle Transaction Manager Server (TMS), which handles object state data automatically. Using any XA Resource Manager imposes specific requirements on how different objects managed by the server application may read and write data to that database, including the following:

- Some XA Resource Managers, such as Oracle7, require that all database operations be scoped within a transaction. This means that the `CourseSynopsisEnumerator` object needs to be scoped within a transaction because this object reads from a database.
- When a transaction is committed or rolled back, the XA Resource Manager automatically handles the durable state implied by the commit or rollback. That is, if the transaction is committed, the XA Resource Manager ensures that all database updates are made permanent. Likewise, if there is a rollback, the XA Resource Manager automatically restores the database to its state prior to the beginning of the transaction.

This characteristic of XA Resource Managers actually makes the design problems associated with handling object state data in the event of a rollback much simpler. Transactional objects can always delegate the commit and rollback responsibilities to the XA Resource Manager, which greatly simplifies the task of implementing a server application.



## Configuration Requirements for the Transactions Sample Application

The University sample applications use an Oracle Transaction Manager Server (TMS). To use the Oracle database, you must include specific Oracle-provided files in the server application build process. For more information about building, configuring, and running the Transactions sample application, see the *Bankapp Sample Using XA* in the WebLogic Enterprise online documentation. For more information about the configurable settings in the `UBBCONFIG` file, see “Modifying the `UBBCONFIG` File to Accommodate Transactions” on page 8-2.



# 4 Transactions in CORBA Client Applications

This topic includes the following sections:

- Overview of WebLogic Enterprise CORBA Transactions
- Summary of the Development Process for Transactions
- Step 1: Use the Bootstrap Object to Obtain the TransactionCurrent Object
- Step 2: Use the TransactionCurrent Methods

This topic describes how to use transactions in CORBA C++, CORBA Java, and ActiveX client applications for the WebLogic Enterprise software. Before you begin, you should read Chapter 1, “Introducing Transactions.”

For an example of how transactions are implemented in working client applications, see the *Bankapp Sample Using XA* in the WebLogic Enterprise online documentation. For an overview of the TransactionCurrent object, see “Client Application Development Concepts” in *Creating CORBA Client Applications*.

# Overview of WebLogic Enterprise CORBA Transactions

*Client applications* use transaction processing to ensure that data remains correct, consistent, and persistent. The transactions in the WebLogic Enterprise software allow client applications to begin and terminate transactions and to get the status of transactions. The WebLogic Enterprise software uses transactions as defined in the *CORBA services Object Transaction Service*, with extensions for ease of use.

Transactions are defined on *interfaces*. The application designer decides which interfaces within a WebLogic Enterprise client/server application will handle transactions. Transaction policies are defined in the Implementation Configuration File (ICF) for C++ server applications, or in the Server Description file (XML) for Java server applications. Generally, the ICF file or the Server Description file for the available interfaces is provided to the client programmer by the application designer.

If you prefer, you can use the Transaction application programming interface (API) defined in the `javax.transaction` package that is shipped with the WebLogic Enterprise (Java) software.

## Summary of the Development Process for Transactions

To add transactions to a client application, complete the following steps:

- Step 1: Use the Bootstrap Object to Obtain the TransactionCurrent Object
- Step 2: Use the TransactionCurrent Methods

The rest of this topic describes these steps using portions of the client applications in the Transactions University sample application. For information about the Transactions University sample application, see the *Bankapp Sample Using XA* in the WebLogic Enterprise online documentation.

The Transactions University sample application is located in the following directory on the WebLogic Enterprise software kit:

- For Microsoft Windows NT systems:  
`drive:\wledir\samples\corba\university\transactions`
- For UNIX systems:  
`drive:/wledir/samples/corba/university/transactions`

# Step 1: Use the Bootstrap Object to Obtain the TransactionCurrent Object

Use the `Bootstrap` object to obtain an object reference to the `TransactionCurrent` object for the specified WebLogic Enterprise domain. For more information about the `TransactionCurrent` object, see “Client Application Development Concepts” in *Creating CORBA Client Applications*.

The following C++, Java, and Visual Basic examples illustrate how the `Bootstrap` object is used to return the `TransactionCurrent` object.

## C++ Example

```
CORBA::Object_var var_transaction_current_oref =  
    Bootstrap.resolve_initial_references("TransactionCurrent");  
CosTransactions::Current_var transaction_current_oref=  
    CosTransactions::Current::_narrow(  
        var_transaction_current_oref.in());
```

## Java Example

```
org.omg.CORBA.Object transCurObj =  
    gBootstrapObjRef.resolve_initial_references(  
        "TransactionCurrent");
```

```
org.omg.CosTransactions.Current gTransCur=  
    org.omg.CosTransactions.CurrentHelper.narrow(transCurObj);
```

### Visual Basic Example

```
Set objTransactionCurrent =  
    objBootstrap.CreateObject("Tobj.TransactionCurrent")
```

## Step 2: Use the TransactionCurrent Methods

The `TransactionCurrent` object has *methods* that allow a client application to manage transactions. These methods can be used to begin and end transactions and to obtain information about the current transaction.

**Note:** Alternatively, a CORBA Java client could use the `UserTransaction` object instead.

Table 4-1 describes the methods in the `TransactionCurrent` object.

**Table 4-1 Methods in the TransactionCurrent Object**

Method	Description
<code>begin</code>	Creates a new transaction. Future operations take place within the scope of this transaction. When a client application begins a transaction, the default transaction timeout is 300 seconds. You can change this default, using the <code>set_timeout</code> method.
<code>commit</code>	Ends the transaction successfully. Indicates that all operations on this client application have completed successfully.
<code>rollback</code>	Forces the transaction to roll back.
<code>rollback_only</code>	Marks the transaction so that the only possible action is to roll back. Generally, this method is used only in server applications.

**Table 4-1 Methods in the TransactionCurrent Object (Continued)**

Method	Description
<code>suspend</code>	Suspends participation in the current transaction. This method returns an object that identifies the transaction and allows the client application to resume the transaction later.
<code>resume</code>	Resumes participation in the specified transaction.
<code>get_status</code>	Returns the status of a transaction with a client application.
<code>get_transaction_name</code>	Returns a printable string describing the transaction.
<code>set_timeout</code>	Modifies the timeout period associated with transactions. The default transaction timeout value is 300 seconds. If a transaction is automatically started instead of explicitly started with the <code>begin</code> method, the timeout value is determined by the value of the <code>TRANTIME</code> parameter in the <code>UBBCONFIG</code> file. For more information about setting the <code>TRANTIME</code> parameter, see Chapter 8, “Administering Transactions.”
<code>get_control</code>	Returns a control object that represents the transaction.

A basic transaction works in the following way:

1. A client application begins a transaction using the `Tobj::TransactionCurrent::begin` method. This method does not return a value.
2. The operations on the CORBA interface execute within the scope of a transaction. If a call to any of these operations raises an exception (either explicitly or as a result of a communications failure), the exception can be caught and the transaction can be rolled back.
3. Use the `Tobj::TransactionCurrent::commit` method to commit the current transaction. This method ends the transaction and starts the processing of the operation. The transaction is committed only if all of the participants in the transaction agree to commit.

The association between the transaction and the client application ends when the client application calls the `Tobj::TransactionCurrent::commit` method or the

Tobj::TransactionCurrent::rollback method. The following C++, Java, and Visual Basic examples illustrate using a transaction to encapsulate the operation of a student registering for a class:

### C++ Example

```
//Begin the transaction
transaction_current_oref->begin();
try {
    //Perform the operation inside the transaction
    pointer_Registar_ref->register_for_courses(student_id, course_number_list);
    ...
    //If operation executes with no errors, commit the transaction:
    CORBA::Boolean report_heuristics = CORBA_TRUE;
    transaction_current_ref->commit(report_heuristics);
}
catch (CORBA::Exception &) {
    //If the operation has problems executing, rollback the
    //transaction. Then throw the original exception again.
    //If the rollback fails, ignore the exception and throw the
    //original exception again.
    try {
        transaction_current_ref->rollback();
    }
    catch (CORBA::Exception &) {
        TP::userlog("rollback failed");
    }

    throw;
}
```

### Java Example

```
try{
    gTransCur.begin();
    //Perform the operation inside the transaction
    not_registered =
        gRegistrarObjRef.register_for_courses(student_id,selected_course_numbers);

    if (not_registered != null)

        //If operation executes with no errors, commit the transaction
```



```
        boolean report_heuristics = true;
        gTransCur.commit(report_heuristics);

    } else gTransCur.rollback();

} catch(org.omg.CosTransactions.NoTransaction nte) {
    System.err.println("NoTransaction: " + nte);
    System.exit(1);
} catch(org.omg.CosTransactions.SubtransactionsUnavailable e) {
    System.err.println("Subtransactions Unavailable: " + e);
    System.exit(1);
} catch(org.omg.CosTransactions.HeuristicHazard e) {
    System.err.println("HeuristicHazard: " + e);
    System.exit(1);
} catch(org.omg.CosTransactions.HeuristicMixed e) {
    System.err.println("HeuristicMixed: " + e);
    System.exit(1);
}
```

## Visual Basic Example

```
' Begin the transaction
'
objTransactionCurrent.begin
'
' Try to register for courses
'
NotRegisteredList = objRegistrar.register_for_courses(mStudentID,
    CourseList, exception)
'
If exception.EX_majorCode = NO_EXCEPTION then
    ' Request succeeded, commit the transaction
    '
    Dim report_heuristics As Boolean
    report_heuristics = True
    objTransactionCurrent.commit report_heuristics
Else
    ' Request failed, Roll back the transaction
    '
    objTransactionCurrent.rollback
    MsgBox "Transaction Rolled Back"
End If
```



# 5 Transactions in EJB Applications

This topic includes the following sections:

- Before You Begin
- General Guidelines
- Transaction Attributes
- Participating in a Transaction
- Transaction Semantics
- Session Synchronization
- Setting Transaction Timeouts

This topic describes how to integrate transactions in Enterprise JavaBeans (EJBs) applications that run under BEA WebLogic Enterprise. Before you begin, you should read Chapter 1, “Introducing Transactions.”

# Before You Begin

Before you begin, you should read Chapter 1, “Introducing Transactions,” particularly the following topics:

- “Transactions in WebLogic Enterprise EJB Applications” on page 1-8
- “Transactions Sample EJB Code” on page 1-25

This document describes the BEA implementation of transactions in Enterprise JavaBeans. The information in this document supplements the Enterprise JavaBeans Specification 1.1, published by Sun Microsystems, Inc.

**Note:** Before proceeding with the rest of this chapter, you must be thoroughly familiar with the *entire* contents of the EJB Specification 1.1 document, particularly the concepts and material presented in Chapter 11, “Support for Transactions.”

For general information about implementing Enterprise JavaBeans in WebLogic Enterprise applications, see “Developing WebLogic Enterprise EJB Applications” in *Getting Started*.

# General Guidelines

The following general guidelines apply when implementing transactions in EJB applications for WebLogic Enterprise:

- WebLogic Enterprise fully supports the EJB Specification 1.1. EJB applications must comply fully with this specification, including all of the various rules, requirements, and limitations that apply to entity beans, stateful session beans, and stateless session beans.
- The EJB specification allows for flat transactions only. Transactions cannot be nested.
- The EJB specification allows for distributed transactions that span multiple resources (such as databases) and supports the two-phase commit protocol. For

more information, see Chapter 7, “Transactions and the WebLogic Enterprise JDBC/XA Driver.”

- For EJB applications running under WebLogic Enterprise, the `AUTOTRAN` setting (if specified) in the `INTERFACES` section of the `UBBCONFIG` file is ignored.
- Use standard programming techniques to optimize transaction processing. For example, properly demarcate transaction boundaries and complete transactions quickly.
- Be sure to tune the EJB cache to ensure maximum performance in transactional EJB applications. For more information, see “Scaling EJB Applications” in *Scaling, Distributing, and Tuning Applications*.

For general guidelines about the WebLogic Enterprise Transaction Service, see “Capabilities and Limitations” on page 2-2.

## Transaction Attributes

This topic includes the following sections:

- About Transaction Attributes for EJBs
- Transaction Attributes for Container-managed Transactions
- Transaction Attributes for Bean-managed Transactions

### About Transaction Attributes for EJBs

Transaction attributes determine how transactions are managed in EJB applications. For each EJB, the transaction attribute specifies whether transactions are demarcated by the WebLogic Enterprise EJB container (container-managed transactions) or by the EJB itself (bean-managed transactions). The setting of the `transaction-type` element in the deployment descriptor determines whether an EJB is container-managed or bean-managed. See Chapter 16, “Deployment Descriptor,” in the EJB Specification 1.1, for more information about the `transaction-type` element.

In general, the use of container-managed transactions is preferred over bean-managed transactions because application coding is simpler. For example, in container-managed transactions, transactions do not need to be started explicitly.

WebLogic Enterprise fully supports method-level transaction attributes as defined in Section 11.4.1 in the EJB Specification 1.1.

# Transaction Attributes for Container-managed Transactions

For container-managed transactions, the transaction attribute is specified in the `container-transaction` element in the deployment descriptor. Container-managed transactions include all Entity beans and any stateful or stateless session beans with a `transaction-type` set to `Container`. For more information about these elements, see Chapter 16, “Deployment Descriptor,” in the EJB Specification 1.1.

The Application Assembler can specify the following transaction attributes for EJBs and their business methods:

- `NotSupported`
- `Supports`
- `Required`
- `RequiresNew`
- `Mandatory`
- `Never`

For a detailed explanation about how the WebLogic Enterprise EJB container responds to these `trans-attribute` settings, see section 11.6.2 in the EJB Specification 1.1.

For EJBs with container-managed transactions, the EJBs have no access to the `javax.transaction.UserTransaction` interface, and the entering and exiting transaction contexts must match. In addition, EJBs with container-managed transactions have limited support for the `setRollbackOnly` and `getRollbackOnly` methods of the `javax.ejb.EJBContext` interface, where invocations are restricted by rules specified in the EJB Specification 1.1.

## Transaction Attributes for Bean-managed Transactions

For bean-managed transactions, the bean specifies transaction demarcations using methods in the `javax.transaction.UserTransaction` interface. Bean-managed transactions include any stateful or stateless session beans with a `transaction-type` set to `Bean`. Entity beans cannot use bean-managed transactions.

For stateless session beans, the entering and exiting transaction contexts must match. For stateful session beans, the entering and exiting transaction contexts may or may not match. If they do not match, the WebLogic Enterprise EJB container maintains associations between the bean and the nonterminated transaction.

Session beans with bean-managed transactions cannot use the `setRollbackOnly` and `getRollbackOnly` methods of the `javax.ejb.EJBContext` interface.

## Participating in a Transaction

When the EJB Specification 1.1 uses the phrase “participating in a transaction,” BEA interprets this to mean that the bean meets either of the following conditions:

- The bean is invoked in a transactional context (container-managed transaction).
- The bean begins a transaction using the `UserTransaction` API in a bean method invoked by the client (bean-managed transaction), and it does *not* suspend or terminate that transaction upon completion of the corresponding bean method invoked by the client.

## Transaction Semantics

This topic contains the following sections:

- Transaction Semantics for Container-managed Transactions
- Transaction Semantics for Bean-managed Transactions

The EJB Specification 1.1 describes semantics that govern transaction processing behavior based on the EJB type (entity bean, stateless session bean, or stateful session bean) and the transaction type (container-managed or bean-managed). These semantics describe the transaction context at the time a method is invoked and define whether the EJB can access methods in the `javax.transaction.UserTransaction` interface. EJB applications must be designed with these semantics in mind.

## Transaction Semantics for Container-managed Transactions

For container-managed transactions, transaction semantics vary for each bean type.

### Transaction Semantics for Stateful Session Beans

Table 5-1 describes the transaction semantics for stateful session beans in container-managed transactions.

**Table 5-1 Transaction Semantics for Stateful Session Beans in Container-managed Transactions**

Method	Transaction Context at the Time the Method Was Invoked	Can Access UserTransaction Methods?
Constructor	Unspecified	No
<code>setSessionContext()</code>	Unspecified	No
<code>ejbCreate()</code>	Unspecified	No
<code>ejbRemove()</code>	Unspecified	No
<code>ejbActivate()</code>	Unspecified	No
<code>ejbPassivate()</code>	Unspecified	No
Business method	Yes or No based on transaction attribute	No
<code>afterBegin()</code>	Yes	No



**Table 5-1 Transaction Semantics for Stateful Session Beans in Container-managed Transactions (Continued)**

Method	Transaction Context at the Time the Method Was Invoked	Can Access UserTransaction Methods?
<code>beforeCompletion()</code>	Yes	No
<code>afterCompletion()</code>	No	No

## Transaction Semantics for Stateless Session Beans

Table 5-2 describes the transaction semantics for stateless session beans in container-managed transactions.

**Table 5-2 Transaction Semantics for Stateless Session Beans in Container-managed Transactions**

Method	Transaction Context at the Time the Method Was Invoked	Can Access UserTransaction Methods?
Constructor	Unspecified	No
<code>setSessionContext()</code>	Unspecified	No
<code>ejbCreate()</code>	Unspecified	No
<code>ejbRemove()</code>	Unspecified	No
Business method	Yes or No based on transaction attribute	No

## Transaction Semantics for Entity Beans

Table 5-3 describes the transaction semantics for entity beans in container-managed transactions.

**Table 5-3 Transaction Semantics for Entity Beans in Container-managed Transactions**

Method	Transaction Context at the Time the Method Was Invoked	Can Access UserTransaction Methods?
Constructor	Unspecified	No
<code>setEntityContext()</code>	Unspecified	No
<code>unsetEntityContext()</code>	Unspecified	No
<code>ejbCreate()</code>	Determined by transaction attribute of matching create	No
<code>ejbPostCreate()</code>	Determined by transaction attribute of matching create	No
<code>ejbRemove()</code>	Determined by transaction attribute of matching remove	No
<code>ejbFind()</code>	Determined by transaction attribute of matching find	No
<code>ejbActivate()</code>	Unspecified	No
<code>ejbPassivate()</code>	Unspecified	No
<code>ejbLoad()</code>	Determined by transaction attribute of business method that invoked <code>ejbLoad()</code>	No
<code>ejbStore()</code>	Determined by transaction attribute of business method that invoked <code>ejbStore()</code>	No
Business method	Yes or No based on transaction attribute	No

# Transaction Semantics for Bean-managed Transactions

For bean-managed transactions, the transaction semantics differ between stateful and stateless session beans. For entity beans, transactions are never bean-managed.

## Transaction Semantics for Stateful Session Beans

Table 5-4 describes the transaction semantics for stateful session beans in bean-managed transactions.

**Table 5-4 Transaction Semantics for Stateful Session Beans in Bean-managed Transactions**

Method	Transaction Context at the Time the Method Was Invoked	Can Access UserTransaction Methods?
Constructor	Unspecified	No
<code>setSessionContext()</code>	Unspecified	No
<code>ejbCreate()</code>	Unspecified	Yes
<code>ejbRemove()</code>	Unspecified	Yes
<code>ejbActivate()</code>	Unspecified	Yes
<code>ejbPassivate()</code>	Unspecified	Yes
Business method	Typically, no <i>unless</i> a previous method execution on the bean had completed while in a transaction context	Yes
<code>afterBegin()</code>	Not applicable	Not applicable
<code>beforeCompletion()</code>	Not applicable	Not applicable
<code>afterCompletion()</code>	Not applicable	Not applicable

## Transaction Semantics for Stateless Session Beans

Table 5-5 describes the transaction semantics for stateless session beans in bean-managed transactions.

**Table 5-5 Transaction Semantics for Stateless Session Beans in Bean-managed Transactions**

Method	Transaction Context at the Time the Method Was Invoked	Can Access UserTransaction Methods?
Constructor	Unspecified	No
<code>setSessionContext()</code>	Unspecified	No
<code>ejbCreate()</code>	Unspecified	Yes
<code>ejbRemove()</code>	Unspecified	Yes
Business method	No	Yes

## Session Synchronization

A stateful session bean using container-managed transactions can implement the `javax.ejb.SessionSynchronization` interface to provide transaction synchronization notifications. In addition, all methods on the stateful session bean must support one of the following transaction attributes: `REQUIRES_NEW`, `MANDATORY` or `REQUIRED`. For more information about the `javax.ejb.SessionSynchronization` interface, see Section 6.5.3 in the EJB Specification 1.1.

If a bean implements `SessionSynchronization`, the WebLogic Enterprise EJB container will typically make the following callbacks to the bean during transaction commit time:

- `afterBegin()`
- `beforeCompletion()`
- `afterCompletion()`

# Setting Transaction Timeouts

Bean providers can specify the timeout period for transactions in EJB applications. If the duration of a transaction exceeds the specified timeout setting, then the Transaction Service rolls back the transaction automatically.

Timeouts are specified according to the transaction type:

- **Container-managed transactions.** The Bean Provider configures the `trans-timeout-seconds` XML element in the `weblogic-ejb-extensions.xml` file. For more information, see the *EJB XML Reference*.
- **Bean-managed transactions.** An application calls the `UserTransaction.setTransactionTimeout` method.

## Handling Exceptions in EJB Transactions

WebLogic Enterprise EJB applications need to catch and handle specific exceptions thrown during transactions. For detailed information about handling exceptions, see Chapter 12, “Exception handling,” in the EJB Specification 1.1 published by Sun Microsystems, Inc.

For more information about how exceptions are thrown by business methods in EJB transactions, see the following tables in Section 12.3: Table 8 (for container-managed transactions) and Table 9 (for bean-managed transactions).

For a client’s view of exceptions, see Section 12.4, particularly Section 12.4.1 (application exceptions), Section 12.4.2 (`java.rmi.RemoteException`), Section 12.4.2.1 (`javax.transaction.TransactionRolledBackException`), and Section 12.4.2.2 (`javax.transaction.TransactionRequiredException`).



# 6 Transactions in RMI Applications

This topic includes the following sections:

- Before You Begin
- General Guidelines

This topic describes how to integrate transactions in RMI applications that run under BEA WebLogic Enterprise.

# Before You Begin

Before you begin, you should read Chapter 1, “Introducing Transactions,” particularly the following topics:

- “Transactions in WebLogic Enterprise RMI Applications” on page 1-11
- “Transactions Sample RMI Code” on page 1-27

For more information about RMI applications, see *Using RMI in a WebLogic Enterprise Environment*.

## General Guidelines

The following general guidelines apply when implementing transactions in RMI applications for WebLogic Enterprise:

- WebLogic Enterprise allows for flat transactions only. Transactions cannot be nested.
- For RMI applications running under WebLogic Enterprise, the AUTOTRAN setting (if specified) in the INTERFACES section of the UBBCONFIG file is ignored.
- Use standard programming techniques to optimize transaction processing. For example, properly demarcate transaction boundaries and complete transactions quickly.
- For RMI applications, callback objects are not recommended for use in transactions because they are not subject to WebLogic Enterprise administration. For more information about callback objects, see “Using RMI with Client-side Callbacks” in *Using RMI in a WebLogic Enterprise Environment*.

For general guidelines about the WebLogic Enterprise Transaction Service, see “Capabilities and Limitations” on page 2-2.



# 7 Transactions and the WebLogic Enterprise JDBC/XA Driver

This topic includes the following sections:

- Before You Begin
- About Transactions and the WebLogic Enterprise JDBC/XA Driver
- JDBC Accessibility in Java Methods
- Using the JDBC/XA Driver
- Implementing Distributed Transactions

This topic describes how to integrate transactions with CORBA Java, EJB, and RMI applications that use the WebLogic Enterprise JDBC/XA driver and run under BEA WebLogic Enterprise. Before you begin, you should read Chapter 1, “Introducing Transactions.”

## Before You Begin

This chapter describes handling transactions in CORBA Java, EJB, and RMI applications that use the WebLogic Enterprise JDBC/XA driver to connect to resources.

For EJB applications, the information in this document supplements the Enterprise JavaBeans Specification 1.1 published by Sun Microsystems, Inc. For general information about implementing Enterprise JavaBeans in WebLogic Enterprise applications, see “Developing WebLogic Enterprise EJB Applications” in *Getting Started*.

## About Transactions and the WebLogic Enterprise JDBC/XA Driver

This topic includes the following sections:

- Support for Transactions Using the WebLogic Enterprise JDBC/XA Driver
- Local Versus Distributed (Global) Transactions
- Transaction Contexts in WebLogic Enterprise JDBC/XA Connections

## Support for Transactions Using the WebLogic Enterprise JDBC/XA Driver

WebLogic Enterprise provides a multithreaded JDBC/XA driver for Oracle Corporation’s Oracle8i database management system. The WebLogic Enterprise JDBC/XA driver fully supports XA, the bidirectional system-level interface between a transaction manager and a resource manager of the X/Open Distributed Transaction Processing (DTP) model. This driver is available to CORBA Java, EJB, and RMI applications and runs in the WebLogic Enterprise environment only.

## Pooled Connections

Java applications use the WebLogic Enterprise JDBC/XA driver to establish concurrent connections to multiple Oracle8i databases via their associated resource managers. For distributed transactions, applications must obtain database connections from the JDBC connection pool. (However, this is not a requirement for other jdbcKona drivers in local transaction mode or for third-party drivers.) Thereafter, applications perform database operations using standard JDBC API calls.

A JDBC connection is governed by the pooled connection lifecycle in the JDBC connection pool. As such, the application server might implicitly close JDBC/XA connections to enforce certain personality-specific transactional resource restrictions, as described in “JDBC Accessibility in Java Methods” on page 7-8. For more information about using WebLogic Enterprise JDBC connection pools with WebLogic Enterprise JDBC/XA driver, see “Using JDBC Connection Pooling” in *Using the JDBC Drivers*.

## Characteristics of JavaServerXA

The JavaServerXA server hosts the WebLogic Enterprise JDBC/XA driver. The JavaServerXA has the following characteristics:

- JavaServerXA is truly multithreaded.
- Multithreaded JavaServerXA cannot use JNI to make database access calls. If an application intends to use JNI to make database access calls, JavaServerXA must be configured to be single-threaded.
- JavaServerXA is still subject to other general multithreaded Java server constraints, as described in “Configuring Multithreaded Java Servers” in *Tuning and Scaling Applications*.
- Each JavaServerXA application can host the WebLogic Enterprise JDBC connection pools that connect to one resource manager only (the resource manager of the Tuxedo group).

### Supported JDBC Standards

WebLogic Enterprise fully supports the JDBC 1.22 API (core functionality), the JDBC 2.0 Core API, and the distributed transactions (the `javax.sql.DataSource` API), connection pooling, and JNDI capabilities in the JDBC 2.0 Optional Package API. See *Using the JDBC Drivers* for a complete list of WebLogic Enterprise-supported JDBC 2.0 features.

## Local Versus Distributed (Global) Transactions

WebLogic Enterprise applications using the WebLogic Enterprise JDBC/XA driver can perform local transactions as well as distributed (also called global) transactions. A local transaction involves updates to a single resource manager (such as a database), while a distributed transaction involves updates across multiple resource managers.

The WebLogic Enterprise JDBC/XA driver never starts a local transaction on behalf of an application. However, if the application performs database operations without first explicitly starting a distributed transaction, then these database operations occur within an “unspecified transaction context” and WebLogic Enterprise delegates the responsibility of handling this situation to the database.

In Oracle8i, for example, the database might start a local transaction to perform such database operations.

- If autocommit is disabled, then it is the application's responsibility to explicitly complete the local transaction by calling the `javax.sql.Connection.commit` or `javax.sql.Connection.rollback` methods.
- If autocommit is enabled, then operations are committed automatically.

Failure to commit a local transaction may result in `XAER_OUTSIDE` error (indicating that the resource manager is performing work outside a distributed transaction) on subsequent distributed transaction operations, which includes beginning a distributed transaction. It is the responsibility of the application to be aware of the transaction context at any point and to complete distributed or local transactions appropriately.

## Differences Between Local and Distributed Transactions

Table 7-1 lists differences between local and distributed transactions.

**Table 7-1 Differences Between Local and Distributed Transactions**

Category	Local Transactions	Distributed Transactions
Resource Managers/Databases	Single database / resource manager	Can span across multiple resource managers
Transaction Demarcation API	Can use the following API: <code>java.sql.Connection</code>	Can use either of following APIs: CORBA API <code>org.omg.CosTransaction</code> <code>TransactionCurrent</code> API EJB API: <code>javax.transaction</code> <code>UserTransaction</code> API
Autocommit	Can be enabled or disabled	Must be disabled

## Configuring the ENABLEXA Parameter in the UBBCONFIG

To use the WebLogic Enterprise JDBC/XA driver, you must specify the `ENABLEXA` parameter (`ENABLEXA=Y`) in the `JDBCONNPOOLS` section of the `UBBCONFIG`, as shown in Listing 7-1. In this example, distributed transactions are *enabled* for the `bank_pool` connection pool.

**Note:** This setting applies only to the WebLogic Enterprise JDBC/XA driver.

**Listing 7-1 Specifying JDBCONNPOOLS Information in UBBCONFIG**

```
JDBCONNPOOLS
  bank_pool
    SRVGRP      = BANK_GROUP1
    SRVID       = 2
    DRIVER      = "weblogic.jdbc20.oci815.Driver"
    URL         = "jdbc:weblogic:oracle:Beq-local"
    PROPS       = "user=scott;password=tiger;server=Beq-Local"
    ENABLEXA    = Y
    INITCAPACITY = 2
    MAXCAPACITY = 10
```

```
CAPACITYINCR      = 1
CREATEONSTARTUP   = Y
```

---

For more information about configuring JDBC connection pools, see “Using JDBC Connection Pooling” in *Using the JDBC Drivers*.

### Demarcating Transaction Boundaries for Local and Distributed Transaction Contexts

Applications must carefully and explicitly demarcate transaction boundaries between distributed and local transaction contexts. For example, when an application uses the WebLogic Enterprise JDBC/XA driver to connect to a database:

- By default, the autocommit feature is automatically disabled because it is assumed that transactions will be distributed.
- For that application to perform local transactions with autocommit (*after* completing the distributed transaction), it must explicitly enable autocommit by calling `javax.sql.Connection.setAutoCommit(true)`.

After completing local transactions, the application must then disable autocommit *before* beginning a new distributed transaction. Listing 7-2 provides a simple example to illustrate switching between a distributed and local transaction.

#### Listing 7-2 Switching Between Distributed and Local Transactions

---

```
// Assumes that javax.transaction.UserTransaction (tx) and
// java.sql.Connection (con) were initialized previously

// Begin a distributed transaction
System.out.println("Beginning distributed transaction...");
tx.begin();
// Database operations within scope of transaction tx
if(gotException){
    try{
        tx.rollback();
        System.out.println("rolled back transaction");
    }catch(Exception e){}
}
elseif{
```

```
        tx.commit();
        System.out.println("committed transaction");
    }
    // Local transactions
    conn.setAutoCommit(true)
    ...[Database operations]...
    conn.setAutoCommit(false)
    // Begin another distributed transaction
    System.out.println("Beginning distributed transaction...");
    tx.begin();
    ...
```

---

## Transaction Contexts in WebLogic Enterprise JDBC/XA Connections

For WebLogic Enterprise JDBC/XA connections, database operations will always be performed in the current transaction context. For example, an application might obtain a JDBC/XA connection in a `NULL` transaction context, begin a distributed transaction, and then perform database operations using that connection. These database operations will be performed in the context of the current distributed transaction.

Applications use WebLogic Enterprise JDBC/XA connection API in the same way as other jdbcKona connections except that, while *within* a distributed transaction context:

- Attempting to enable autocommit mode by calling the `javax.sql.Connection.setAutoCommit` method on the WebLogic Enterprise JDBC/XA connection will throw a `SQLException`.
- Attempting to complete the distributed transaction by calling the `javax.sql.Connection.commit` or `javax.sql.Connection.commit` methods on the WebLogic Enterprise JDBC/XA connection will throw a `SQLException`.

Listing 7-3 shows, in a sample CORBA Java application, how to determine the current transaction context and commit a local or global transaction accordingly.

### Listing 7-3 Determining Whether the Application Is in a Distributed

### Transaction

---

```
// Assumes that org.omg.CosTransactions.Current (tc) and
// java.sql.Connection (con) were initialized before
// database operations were attempted
if (tc.get_status() !=
org.omg.CosTransactions.Status.StatusNoTransaction)
{
    // Application is currently in a distributed transaction
    tc.commit(true);
}
else
{
    // Application is currently in a local transaction
    con.commit();
}
```

---

Similarly, for bean-managed transactions in an EJB application, the application can determine whether the application is currently in a distributed transaction by calling the `UserTransaction.getStatus()` method and testing for a returned `STATUS_NO_TRANSACTION`.

## JDBC Accessibility in Java Methods

This topic includes the following sections:

- JDBC/XA Accessibility in CORBA Methods
- JDBC/XA Accessibility in EJB Methods

**Note:** Attempting to use a WebLogic Enterprise JDBC/XA connection in a method where it is not supported may have undefined behavior and possibly raise a `SQLException`.



## JDBC/XA Accessibility in CORBA Methods

Table 7-2 lists which methods in CORBA methods can access JDBC/XA connections.

**Table 7-2 JDBC/XA Connection Accessibility for CORBA Objects**

Server Method	Accessibility
Constructor	Not supported
<code>initialize</code>	Supported, after <code>open_xa_rm</code>
<code>activate_obj</code>	Supported
<code>deactivate_obj</code>	Supported
Business method	Supported
<code>release</code>	Supported, before <code>close_xa_rm</code>

After completing the `initialize` method, WebLogic Enterprise automatically closes any open connections and writes a warning message to the `ULOG`.

For transaction-bound and process-bound objects, the CORBA framework allows open connections to be retained at method end, and the transaction context of the retained connections will be as described in “Transaction Contexts in WebLogic Enterprise JDBC/XA Connections” on page 7-7 upon subsequent method invocations. However, for method-bound objects, applications *must* explicitly close open connections before method end. If not, WebLogic Enterprise automatically closes any open connections and writes a warning message to the `ULOG`.

## JDBC/XA Accessibility in EJB Methods

For EJB methods, accessibility to JDBC/XA connections varies depending on the EJB type. For details about retaining JDBC/XA connections across method invocations (for stateful session beans only), including examples, see Section 11.3.3 in the Enterprise JavaBeans Specification 1.1, published by Sun Microsystems, Inc.

**Note:** For all bean types, after completing the `ejbCreate` method, WebLogic Enterprise automatically closes any open connections and writes a warning message to the `ULOG`.

### Stateful Session Beans

Table 7-3 lists which stateful session bean methods can access JDBC/XA connections.

**Table 7-3 JDBC/XA Connection Accessibility for Stateful Session Beans**

Bean Method	Container-managed Transaction	Bean-managed Transaction
Constructor	Not supported	Not supported
<code>setSessionContext</code>	Not supported	Not supported
<code>ejbCreate</code> <code>ejbRemove</code> <code>ejbActivate</code> <code>ejbPassivate</code>	Supported, but in unspecified transaction context (as defined in the Enterprise JavaBeans 1.1 specification)	Supported, but in unspecified transaction context (as defined in the Enterprise JavaBeans 1.1 specification), unless the bean explicitly begins a transaction using <code>UserTransaction</code>
Business method	Supported	Supported
<code>afterBegin</code>	Supported	N/A
<code>beforeCompletion</code>	Supported	N/A
<code>afterCompletion</code>	Supported	N/A

For stateful session beans, the Bean Provider must close all JDBC connections in `ejbPassivate` and assign the instance's fields storing the connections to null. However, after completing the `ejbPassivate` method, WebLogic Enterprise automatically closes any open connections and writes a warning message to the `ULOG`.

## Stateless Session Beans

Table 7-4 lists which stateless session bean methods can access JDBC/XA connections.

**Table 7-4 JDBC/XA Connection Accessibility for Stateless Session Beans**

<b>Bean Method</b>	<b>Container-managed Transaction</b>	<b>Bean-managed Transaction</b>
Constructor	Not supported	Not supported
<code>setSessionContext</code>	Not supported	Not supported
<code>ejbCreate</code>	Not supported	Not supported
<code>ejbRemove</code>	Not supported	Not supported
Business method	Supported	Supported

**Note:** For stateless session beans, after completing a business method, WebLogic Enterprise automatically closes any open connections and writes a warning message to the `ULOG`.

## Entity Beans

Table 7-5 lists which entity bean methods can access JDBC/XA connections.

**Table 7-5 JDBC/XA Connection Accessibility for Entity Beans**

<b>Bean Method</b>	<b>Accessibility</b>
Constructor	Not supported
<code>setEntityContext</code>	Not supported
<code>unsetEntityContext</code>	Not supported
<code>ejbCreate</code>	Supported
<code>ejbPostCreate</code>	Supported
<code>ejbRemove</code>	Supported

**Table 7-5 JDBC/XA Connection Accessibility for Entity Beans (Continued)**

Bean Method	Accessibility
ejbFind	Supported
ejbActivate	Not supported
ejbPassivate	Not supported
ejbLoad	Supported
ejbStore	Supported
business method	Supported

## Using the JDBC/XA Driver

Before applications can use the WebLogic Enterprise JDBC/XA driver, the JDBC/XA driver must be integrated into your development environment by completing the following steps:

1. Build the multithreaded `JavaServerXA` application, binding it with the Oracle8i Resource Manager, as described in “Using the WebLogic Enterprise JDBC/XA Driver” in *Using the JDBC Drivers*.
2. In the `UBBCONFIG`, configure the `OPENINFO` parameter in the `GROUPS` section according to the definition of the `XA` parameter for the Oracle database.  
Listing 7-4 shows an example of an `OPENINFO` setting in a sample `UBBCONFIG`.

---

**Listing 7-4 OPENINFO Setting in Sample UBBCONFIG**

---

```
*GROUPS
  SYS_GRP
    LMID      = SITE1
    GRPNO     = 1
  BANK_GROUP1
    LMID      = SITE1
    GRPNO     = 2
  OPENINFO =
```

```
"ORACLE_XA:Oracle_XA+Acc=P/scott/tiger+SesTm=100+LogDir=..+DbgFl=0
x7+MaxCur=15+Threads=true"
TMSNAME   = TMS_ORA
TMSCOUNT = 2
```

---

For more information about the XA parameter, see the “A Oracle XA” chapter in the Fundamentals section of the Oracle Corporation *Oracle8i Application Developer's Guide*.

3. If you want the JavaServerXA to be multithreaded, you must specify the -M option for the CLOPT parameter, which is defined in the JavaServerXA entry in the SERVERS section of the UBBCONFIG file.

**Note:** For single-threaded JavaServerXA operation, skip this step.

Listing 7-5 shows an example of JavaServerXA configured for multithreading in a sample UBBCONFIG

#### **Listing 7-5 Multithreaded Server Configuration in Sample UBBCONFIG**

---

```
*SERVERS
  DEFAULT:
    RESTART = Y
    MAXGEN  = 5
    ...
  JavaServerXA
    SRVGRP  = BANK_GROUP1
    SRVID   = 2
    SRVTYPE = JAVA
    CLOPT   = "-A -- -M 10 BankApp.jar TellerFactory_1 bank_pool"
    RESTART = N
```

---

To specify connection pooling, you need to specify SRVTYPE=JAVA in the SERVERS section.

4. In the `UBBCONFIG`, configure the WebLogic Enterprise JDBC/XA driver in the WebLogic Enterprise JDBC Connection Pool, as described in “Using the WebLogic Enterprise JDBC/XA Driver” in *Using the JDBC Drivers*. Listing 7-6 shows an example of JDBC connection pool settings for a connection pool named `bank_pool` in a sample `UBBCONFIG`.

### Listing 7-6 JDBC Connection Pool Settings in Sample UBBCONFIG

---

```
*JDBCONNPOOLS
  bank_pool
    SRVGRP      = BANK_GROUP1
    SRVID       = 2
    DRIVER      = "weblogic.jdbc20.oci815.Driver"
    URL         = "jdbc:weblogic:oracle:beq-local"
    PROPS       = "user=scott;password=tiger;server=Beq-Local"
    ENABLEXA    = Y
    INITCAPACITY = 2
    MAXCAPACITY = 10
    CAPACITYINCR = 1
    CREATEONSTARTUP = Y
```

---

5. Boot the `JavaServerXA` application, as described in “Using the WebLogic Enterprise JDBC/XA Driver” in *Using the JDBC Drivers*.

# Implementing Distributed Transactions

This topic includes the following sections:

- Importing Packages
- Initializing the `TransactionCurrent` Object Reference
- Finding the Connection Pool via JNDI
- Setting Up XA Distributed Transactions
- Performing a Distributed Transaction

In addition to the fully supported examples supplied on the CD-ROM with this release of WebLogic Enterprise, the BEA WebLogic Enterprise team provides several unsupported code examples on a password protected Web site for WebLogic Enterprise customers. The code samples in this topic come from a version of the WebLogic Enterprise XA Bankapp sample application that is available from the unsupported samples WebLogic Enterprise Web site. The URL for the unsupported samples WebLogic Enterprise Web site is specified in the product Release Notes under “About This BEA WebLogic Enterprise Release” in the subsection “Unsupported Samples and Tools Web Page.”

This application is different from the one described in the *Bankapp Sample Using XA* in the WebLogic Enterprise online documentation.

**Note:** This topic does not attempt to fully describe this sample application. It merely uses code fragments to illustrate the use of the JDBC/XA driver in a CORBA application.

## Importing Packages

Listing 7-7 shows the packages that the application imports. In particular, note that:

- The `java.sql.*` and `javax.sql.*` packages are required for database operations.
- The `javax.naming.*` package is required for performing a JNDI lookup on the pool name, which is passed in as a command-line parameter upon server startup. The pool name must be registered on that server group.

### Listing 7-7 Importing Required Packages

---

```
import java.sql.*;
import javax.sql.*;
import javax.naming.*;
import com.beasys.Tobj.*;
```

---

## Initializing the TransactionCurrent Object Reference

Listing 7-8 shows initializing the `TransactionCurrent` object reference, which will be used by the `Teller` operations to start and stop transactions.

### **Listing 7-8    Initializing the TransactionCurrent Object Reference**

---

```
static org.omg.CosTransactions.Current trans_cur_ref;  
  
org.omg.CORBA.Object trans_cur_oref =  
TP.bootstrap().resolve_initial_references("TransactionCurrent");
```

---

## Finding the Connection Pool via JNDI

Listing 7-9 shows finding the connection pool via JNDI. The connection pool name is registered on the server group and is passed in as a command-line parameter upon server startup. Subsequent database connections are obtained from this pool.

### **Listing 7-9    Finding the Connection Pool via JNDI**

---

```
static DataSource pool;  
  
...  
  
public void get_connpool(String pool_name)  
    throws Exception  
    {  
    try {  
        javax.naming.Context ctx = new InitialContext();  
        pool = (DataSource)ctx.lookup("jdbc/" + pool_name);  
    }  
    catch (javax.naming.NamingException ex){  
        TP.userlog("Couldn't obtain JDBC connection pool: " +  
pool_name);  
        throw ex;  
    }  
    }  
}
```



## Setting Up XA Distributed Transactions

Listing 7-10 shows setting up XA distributed transactions by calling the `open_xa_rm` method (in `server.initialize`) and obtaining a reference to the `TransactionCurrent` object.

**Note:** This step is required for CORBA applications but not for EJB or RMI applications.

### Listing 7-10 Setting Up XA Distributed Transactions

---

```
TP.open_xa_rm();

org.omg.CORBA.Object trans_cur_oref =
TP.bootstrap().resolve_initial_references("TransactionCurrent");

trans_cur_ref =
org.omg.CosTransactions.CurrentHelper.narrow(trans_cur_oref);
```

---

## Performing a Distributed Transaction

Listing 7-11 shows a complete distributed transaction that involves the transfer of money from one bank account to another.

### Sequence of Tasks

The application performs the distributed application in the following sequence:

1. The application calls the `begin` method to start the transaction.
2. The application performs the following database operations:
  - Withdrawing the money from one account
  - Depositing the money into another account.

3. The application updates balances.
4. The application catches any exceptions thrown during the database operations.
5. The application closes the distributed transaction and updates teller statistics.
  - If an exception was thrown during the database operations, the application rolls back the transaction by calling the `rollback` method.
  - If no exceptions were thrown, the application commits the transaction by calling the `commit` method.

### Listing 7-11 Performing a Distributed Transaction

---

```
public void transfer(int fromAccountID, int toAccountID, float
amount, BalanceAmountsHolder balances)
    throws AccountRecordNotFound, IOException, InsufficientFunds
{
    boolean success = false;

    try {
        // Increment the number of requests the teller has received.
        tellerStats.totalTellerRequests += 1;

        // Begin the global transaction.
        BankAppServerImpl.trans_cur_ref.begin();

        // Flag this as a transfer.
        transferInProgress = true;

        // Perform the withdrawal first.
        float withdrawalBalance = withdraw(fromAccountID, amount);

        // Perform the deposit next.
        float depositBalance = deposit(toAccountID, amount);

        balances.value = new BalanceAmounts();
        balances.value.fromAccount = withdrawalBalance;
        balances.value.toAccount   = depositBalance;

        success = true;

        // Catch any exceptions thrown during database operations
    }
    catch (AccountRecordNotFound e) {
        throw e;
    }
}
```

```
        catch (InsufficientFunds e) {
            throw e;
        }
        catch (IOException e) {
            throw e;
        }
        catch(Exception e) {
            TP.userlog("Exception caught in transfer(): "
                + e.getMessage());
            e.printStackTrace();
            throw new org.omg.CORBA.INTERNAL();
        }
        finally {
            try {
                // Complete the distributed transaction and
                // update the Teller statistics.
                if (success) {
                    tellerStats.totalTellerSuccess += 1;
                    BankAppServerImpl.trans_cur_ref.commit(true);
                } else {
                    tellerStats.totalTellerFail += 1;
                    BankAppServerImpl.trans_cur_ref.rollback();
                }
            }
            catch(Exception e) {
                TP.userlog("Unexpected Exception thrown during commit or
rollback: " + e.getMessage());
                e.printStackTrace();
                throw new org.omg.CORBA.INTERNAL();
            }
            transferInProgress = false;
        }
    }
}
```

---

## The withdraw Method

Listing 7-12 shows the withdraw method that is invoked in Listing 7-11. The withdraw method shows accessing the database to withdraw money from the specified account.

### Listing 7-12 withdraw method

---

```
public float withdraw(int accountID, float amount)
    throws AccountRecordNotFound,
           IOException,
           InsufficientFunds,
           TellerInsufficientFunds
{
    boolean success = false;

    try {
        if (!transferInProgress) {
            // This is just a plain withdrawal; it is NOT a transfer.

            // Increment the number of requests that this teller
            // has received.
            tellerStats.totalTellerRequests += 1;

            // Decrement the balance left in the Teller's ATM machine.
            tellerStats.totalTellerBalance -= amount;

            // Begin the global transaction.
            BankAppServerImpl.trans_cur_ref.begin();

            // Check to see if the minimum TELLER threshold balance
            // has not been reached; if so, amount will be added back in
            // in the finally clause.
            if (tellerStats.totalTellerBalance < MinTellerBalance)
                throw new TellerInsufficientFunds();
        }

        AccountDataHolder accountDataHolder =
            new AccountDataHolder(new AccountData());
        accountDataHolder.value.accountID = accountID;
        accountDataHolder.value.balance = -amount;

        // Withdraw the money from the account.
        theDBAccess_ref.update_account(accountDataHolder);
        success = true;
        return(accountDataHolder.value.balance);
    }
    catch (AccountRecordNotFound e) {
        throw e;
    }
    catch (InsufficientFunds e) {
        throw e;
    }
    catch (TellerInsufficientFunds e) {
```

```
        throw e;
    }
    catch (DataBaseException e) {
        throw new IOException();
    }
    catch (Exception e) {
        TP.userlog("Exception caught in withdraw(): "
            + e.getMessage());
        e.printStackTrace();
        throw new org.omg.CORBA.INTERNAL();
    }
    finally {
        // Terminate the transaction and update the Teller statistics.
        if (!transferInProgress) {
            try {
                if (success) {
                    tellerStats.totalTellerSuccess += 1;
                    BankAppServerImpl.trans_cur_ref.commit(true);
                } else {
                    tellerStats.totalTellerFail += 1;
                    tellerStats.totalTellerBalance += amount;
                    BankAppServerImpl.trans_cur_ref.rollback();
                }
            }
            catch (Exception e) {
                TP.userlog("Unexpected Exception thrown during commit or
rollback: " + e.getMessage());
                e.printStackTrace();
                throw new org.omg.CORBA.INTERNAL();
            }
        }
    }
}
```

---

## The deposit Method

Listing 7-13 shows the deposit method that is invoked in Listing 7-11. The deposit method shows accessing the database deposit money to the specified account.

### **Listing 7-13** deposit method

---

```
public float deposit(int accountID, float amount)
    throws AccountRecordNotFound, IOException
{
    boolean success = false;
```

```

        try {
            // If this is a transfer request, then the global transaction
            // was started in the TellerImpl.transfer method; otherwise,
            // start the transaction here.
            if (!transferInProgress) {

                // This is just a plain deposit; it is NOT a transfer.
                // Increment the number of requests that this teller
                // has received.
                tellerStats.totalTellerRequests += 1;

                // Begin the global transaction.
                BankAppServerImpl.trans_cur_ref.begin();
            }

            AccountDataHolder accountDataHolder =
            new AccountDataHolder(new AccountData());
            accountDataHolder.value.accountID = accountID;
            accountDataHolder.value.balance = amount;

            // Deposit the money in the account.
            theDBAccess_ref.update_account(accountDataHolder);

            success = true;
            return(accountDataHolder.value.balance);
        }
        catch (AccountRecordNotFound e) {
            throw e;
        }
        catch (DataBaseException e) {
            throw new IOException();
        }
        catch (Exception e) {
            TP.userlog("Exception caught in BankApp.deposit(): "
            + e.getMessage());
            e.printStackTrace();
            throw new org.omg.CORBA.INTERNAL();
        }
        finally {
            try {
                // Terminate the transaction and update the Teller statistics.
                if (!transferInProgress) {
                    if (success) {
                        tellerStats.totalTellerSuccess += 1;
                        BankAppServerImpl.trans_cur_ref.commit(true);
                    }
                    else {
                        tellerStats.totalTellerFail += 1;
                        BankAppServerImpl.trans_cur_ref.rollback();
                    }
                }
            }
        }
    }

```

```
        }
        catch(Exception e) {
TP.userlog("Unexpected Exception thrown during commit or rollback:
"
        + e.getMessage());
e.printStackTrace();
throw new org.omg.CORBA.INTERNAL();
        }
    }
}
```

---





# 8 Administering Transactions

This topic includes the following sections:

- Modifying the UBBCONFIG File to Accommodate Transactions
- Modifying the Domain Configuration File to Support Transactions (WebLogic Enterprise Servers)
- Sample Distributed Application Using Transactions

Before you begin, you should read Chapter 1, “Introducing Transactions.” In addition, for container-managed transaction demarcation in EJB applications, you can configure the transaction timeout setting, as described in “Setting Transaction Timeouts” on page 5-11.

# Modifying the UBBCONFIG File to Accommodate Transactions

This topic includes the following sections:

- Summary of Steps
- Step 1: Specify Application-wide Transactions in the RESOURCES Section
- Step 2: Create a Transaction Log (TLOG)
- Step 3: Define Each Resource Manager (RM) and the Transaction Manager Server in the GROUPS Section
- Step 4: Enable an Interface to Begin a Transaction

## Summary of Steps

To accommodate transactions, you must modify the RESOURCES, MACHINES, GROUPS, and the INTERFACES or SERVICES sections of the application's UBBCONFIG file in the following ways:

- In the RESOURCES section, specify the application-wide number of allowed transactions and the value of the commit control flag.
- In the MACHINES section, create the TLOG information for each machine.
- In the GROUPS section, indicate information about each resource manager and about the transaction manager server.
- In the INTERFACES section (WebLogic Enterprise System for CORBA applications only) or the SERVICES section (BEA Tuxedo System), enable the automatic transaction option. This option does *not* apply to EJB or RMI applications.

For instructions about modifying these sections in the UBBCONFIG, see “Creating a Configuration File” in the *Administration Guide*.

## Step 1: Specify Application-wide Transactions in the RESOURCES Section

Table 8-1 provides a description of transaction-related parameters in the RESOURCES section of the configuration file.

**Table 8-1 Transaction-Related Parameters in the RESOURCES Section**

Parameter	Meaning
MAXGTT	<p>Limits the total number of global transaction identifiers (GTRIDs) allowed on one machine at one time. The maximum value allowed is 2048, the minimum is 0, and the default is 100. You can override this value on a per-machine basis in the MACHINES section.</p> <p>Entries remain in the table only while the global transaction is active, so this parameter has the effect of setting a limit on the number of simultaneous transactions.</p>
CMTRET	<p>Specifies the initial setting of the TP_COMMIT_CONTROL characteristic. The default is COMPLETE. Following are its two settings:</p> <ul style="list-style-type: none"> <li>■ LOGGED—the TP_COMMIT_CONTROL characteristic is set to TP_CMT_LOGGED, which means that <code>tpcommit ( )</code> returns when all the participants have successfully precommitted.</li> <li>■ COMPLETE—the TP_COMMIT_CONTROL characteristic is set to TP_CMT_COMPLETE, which means that <code>tpcommit ( )</code> will not return until all the participants have successfully committed.</li> </ul> <p><b>Note:</b> You should consult with the RM vendors to determine the appropriate setting. If any RM in the application uses the <i>late commit</i> implementation of the XA standard, the setting should be COMPLETE. If all the resource managers use the <i>early commit</i> implementation, the setting should be LOGGED for performance reasons. (You can override this setting with <code>tpscmt ( . )</code>.)</p>

## Step 2: Create a Transaction Log (TLOG)

This section discusses creating a transaction log (TLOG), which refers to a log in which information on transactions is kept until the transaction is completed.

### Creating the UDL

The Universal Device List (UDL) is like a map of the BEA Tuxedo file system. The UDL gets loaded into shared memory when an application is booted. To create an entry in the UDL for the TLOG device, create the UDL on each machine using global transactions. If the TLOGDEVICE is mirrored between two machines, it is unnecessary to do this on the paired machine. The Bulletin Board Liaison (BBL) then initializes and opens the TLOG during the boot process.

To create the UDL, enter a command using the following format, before the application has been booted:

```
tmadmin -c crdl -z config -b blocks
```

where:

<code>-z config</code>	Specifies the full path name for the device where you should create the UDL.
<code>-b blocks</code>	Specifies the number of blocks to be allocated on the device.
<code>config</code>	Should match the value of the TLOGDEVICE parameter in the MACHINES section of the UBBCONFIG file.

**Note:** In general, the value that you supply for `blocks` should not be less than the value for `TLOGSIZE`. For example, if `TLOGSIZE` is specified as 200 blocks, specifying `-b 500` would not cause a degradation.

For more information about storing the TLOG, see the *Installation Guide*.

### Defining Transaction-related Parameters in the MACHINES Section

You can define a global transaction log (TLOG) using several parameters in the MACHINES section of the UBBCONFIG file. You must manually create the device list entry for the TLOGDEVICE on each machine where a TLOG is needed. You can do this either before or after TUXCONFIG has been loaded, but it must be done before the system is booted.

**Note:** If you are not using transactions, the TLOG parameters are not required.

Table 8-2 provides a description of transaction-related parameters in the MACHINES section of the configuration file.

**Table 8-2 Transaction-related Parameters in the MACHINES Section**

Parameter	Meaning
TLOGNAME	The name of the DTP transaction log for this machine.
TLOGDEVICE	Specifies the WebLogic Enterprise or BEA Tuxedo file system that contains the DTP transaction log (TLOG) for this machine. If this parameter is not specified, the machine is assumed not to have a TLOG. The maximum string value length is 64 characters.
TLOGSIZE	The size of the TLOG file in physical pages. Its value must be between 1 and 2048, and its default is 100. The value should be large enough to hold the number of outstanding transactions on the machine at a given time. One transaction is logged per page. The default should suffice for most applications.
TLOGOFFSET	<p>Specifies the offset in pages from the beginning of TLOGDEVICE to the start of the VTOC that contains the transaction log for this machine. The number must be greater than or equal to 0 and less than the number of pages on the device. The default is 0.</p> <p>TLOGOFFSET is rarely necessary. However, if two VTOCs share the same device or if a VTOC is stored on a device (such as a file system) that is shared with another application, you can use TLOGOFFSET to indicate a starting address relative to the address of the device.</p>

### Creating the Domains Transaction Log (BEA Tuxedo Servers)

This section applies to the BEA Tuxedo system only.

You can create the Domains transaction log before starting the Domains gateway group by using the following command:

```
dmadmin(1) crdmlog (crdlog) -d local_domain_name
```

Create the Domains transaction log for the named local domain on the current machine (the machine on which `dmadmin` is running). The command uses the parameters specified in the `DMCONFIG` file. This command fails if the named local domain is active on the current machine or if the log already exists. If the transaction log has not been created, the Domains gateway group creates the log when it starts up.

## Step 3: Define Each Resource Manager (RM) and the Transaction Manager Server in the GROUPS Section

Additions to the GROUPS section fall into two categories:

- Defining the transaction manager servers that perform most of the work that controls global transactions:
  - The TMSNAME parameter specifies the name of the server executable.
  - The TMSCOUNT parameter specifies the number of such servers to boot (the minimum is 2, the maximum is 10, and the default is 3).

A null transactional manager server does not communicate with any resource manager. It is used to exercise an application's use of the transactional primitives before actually testing the application in a recoverable, *real* environment. This server is named TMS and it simply begins, commits, or terminates without talking to any resource manager.

- Defining opening and closing information for each resource manager:
  - OPENINFO is a string with information used to open a resource manager.
  - CLOSEINFO is used to close a resource manager.

### Sample GROUPS Section

The following sample GROUPS section derives from the bankapp banking application:

```
BANKB1 GRPNO=1 TMSNAME=TMS_SQL TMSCOUNT=2  
OPENINFO="TUXEDO/SQL:<APPDIR>/bankd11:bankdb:readwrite"
```

Table 8-3 describes the transaction values specified in this sample GROUPS section.

**Table 8-3 Transaction Values in the GROUPS Section of a Sample UBBCONFIG File**

Transaction Value	Meaning
BANKB1 GRPNO=1 TMSNAME=TMS_SQL\ TMSCOUNT=2	Contains the name of the transaction manager server (TMS_SQL) and the number (2) of these servers to be booted in the group BANKB1

**Table 8-3 Transaction Values in the GROUPS Section of a Sample UBBCONFIG File (Continued)**

Transaction Value	Meaning
TUXEDO/SQL	Published name of the resource manager
<APPDIR>/bankd11	Includes a device name
bankdb	Database name
readwrite	Access mode

## Characteristics of the TMSNAME, TMSCOUNT, OPENINFO, and CLOSEINFO Parameters

Table 8-4 lists the characteristics of the TMSNAME, TMSCOUNT, OPENINFO, and CLOSEINFO parameters.

**Table 8-4 Characteristics of the TMSNAME, TMSCOUNT, OPENINFO, and CLOSEINFO Parameters**

Parameter	Characteristics
TMSNAME	Name of the transaction manager server executable. Required parameter for transactional configurations. TMS is a null transactional manager server.
TMSCOUNT	Number of transaction manager servers (must be between 2 and 10). Default is 3.
OPENINFO CLOSEINFO	Represents information to open or close a resource manager. Content depends on the specific resource manager. Starts with the name of the resource manager. Omission means the resource manager needs no information to open.

# Step 4: Enable an Interface to Begin a Transaction

To enable an interface to begin a transaction, you change different sections in the `UBBCONFIG` file, depending on whether you are configuring a WebLogic Enterprise CORBA server or BEA Tuxedo server:

- Changing the INTERFACES Section (WebLogic Enterprise CORBA Servers)
- Changing the SERVICES Section (BEA Tuxedo Servers)

## Changing the INTERFACES Section (WebLogic Enterprise CORBA Servers)

The `INTERFACES` section in the `UBBCONFIG` file supports WebLogic Enterprise CORBA interfaces:

- For each CORBA interface, set `AUTOTRAN` to `Y` if you want a transaction to start automatically when an operation invocation is received. `AUTOTRAN=Y` has no effect if the interface is already in transaction mode. The default is `N`. The effect of specifying a value for `AUTOTRAN` depends on the transactional policy specified by the developer in the Implementation Configuration File (ICF) in C++, or the Server Description File (XML) in Java, for the interface. This transactional policy will become the transactional policy attribute of the associated `T_IFQUEUE MIB` object at run time. The only time this value affects the behavior of the application is if the developer specified a transaction policy of `optional`.

**Note:** To work properly, this feature depends on collaboration between the system designer and the administrator. If the administrator sets this value to `Y` without prior knowledge of the transaction policy defined by the developer in the interface's ICF or XML file, the actual run time effect of the parameter might be unknown.

- If `AUTOTRAN` is set to `Y`, you must set the `TRANTIME` parameter, which specifies the transaction timeout, in seconds, for the transactions to be created. The value must be greater than or equal to zero and must not exceed 2,147,483,647 ( $2^{31} - 1$ , or about 70 years). A value of zero implies there is no timeout for the transaction. (The default is 30 seconds.)

**Note:** For EJB and RMI applications, the `AUTOTRAN` and `TRANTIME` settings are ignored.



Table 8-5 describes the characteristics of the AUTOTRAN, TRANTIME, and FACTORYROUTING parameters.

**Table 8-5 Characteristics of the AUTOTRAN, TRANTIME, and FACTORYROUTING Parameters**

Parameter	Characteristics
AUTOTRAN	<ul style="list-style-type: none"><li>■ Makes an interface the initiator of a transaction.</li><li>■ To work properly, it is dependent on collaboration between the system designer and the system administrator. If the administrator sets this value to Y without prior knowledge of the ICF or XML transaction policy set by the developer, the actual run-time effort of the parameter might be unknown.</li><li>■ The only time this value affects the behavior of the application is if the developer specified a transaction policy of optional.</li><li>■ If a transaction already exists, a new one is not started.</li><li>■ Default is N.</li></ul>
TRANTIME	<ul style="list-style-type: none"><li>■ Represents the timeout for the AUTOTRAN transactions.</li><li>■ Valid values are between 0 and <math>2^{31} - 1</math>, inclusive.</li><li>■ Zero (0) represents no timeout.</li><li>■ Default is 30 seconds.</li></ul>
FACTORYROUTING	<ul style="list-style-type: none"><li>■ Specifies the name of the routing criteria to be used for factory-based routing for this CORBA interface.</li><li>■ You must specify a FACTORYROUTING parameter for interfaces requesting factory-based routing.</li></ul>

## Changing the SERVICES Section (BEA Tuxedo Servers)

The following are three transaction-related features in the SERVICES section:

- If you want a service (instead of a client) to begin a transaction, you must set the AUTOTRAN flag to Y. This is useful if the service is not needed as part of any larger transaction, and if the application wants to relieve the client of making transaction decisions. If the service is called when there is already an existing transaction, this call becomes part of it. (The default is N.)

**Note:** Generally, clients are the best initiators of transactions because a service has the potential of participating in a larger transaction.

- If AUTOTRAN is set to Y, you must set the TRANTIME parameter, which is the transaction timeout, in seconds, for the transactions to be created. The value must be greater than or equal to 0 and must not exceed 2,147,483,647 ( $2^{31} - 1$ , or about 70 years). A value of zero implies there is no timeout for the transaction. (The default is 30 seconds.)

**Note:** For EJB and RMI applications, the AUTOTRAN and TRANTIME settings are ignored.

- You must specify a ROUTING parameter for transactions that request data-dependent routing.

Table 8-6 describes the characteristics of the AUTOTRAN, TRANTIME, and ROUTING parameters:

**Table 8-6 Characteristics of the AUTOTRAN, TRANTIME, and ROUTING Parameters**

Parameter	Characteristics
AUTOTRAN	<p>Makes a service the initiator of a transaction.</p> <p>Relieves the client of the transactional burden.</p> <p>If a transaction already exists, a new one is not started.</p> <p>Default is N.</p>
TRANTIME	<p>Represents the timeout for the AUTOTRAN transactions.</p> <p>Valid values are between 0 and <math>2^{31} - 1</math>, inclusive.</p> <p>0 represents no timeout.</p> <p>Default is 30 seconds.</p>
ROUTING	<p>Points to an entry in the ROUTING section where data-dependent routing is specified for transactions that request this service.</p>

# Modifying the Domain Configuration File to Support Transactions (WebLogic Enterprise Servers)

This topic includes the following sections:

- Characteristics of the DMTLOGDEV, DMTLOGNAME, DMTLOGSIZE, MAXRDTRAN, and MAXTRAN Parameters
- Characteristics of the AUTOTRAN and TRANTIME Parameters (WebLogic Enterprise CORBA and Tuxedo Servers)

To enable transactions across domains, you need to set parameters in both the `DM_LOCAL_DOMAINS` and the `DM_REMOTE_SERVICES` sections of the Domains configuration file (`DMCONFIG`). Entries in the `DM_LOCAL_DOMAINS` section define local domain characteristics. Entries in the `DM_REMOTE_SERVICES` section define information on services that are *imported* and that are available on remote domains.

## Characteristics of the DMTLOGDEV, DMTLOGNAME, DMTLOGSIZE, MAXRDTRAN, and MAXTRAN Parameters

The `DM_LOCAL_DOMAINS` section of the Domains configuration file identifies local domains and their associated gateway groups. This section must have an entry for each gateway group (Local Domain). Each entry specifies the parameters required for the Domains gateway processes running in that group.

Table 8-7 provides a description of the five transaction-related parameters in this section: `DMTLOGDEV`, `DMTLOGNAME`, `DMTLOGSIZE`, `MAXRDTRAN`, and `MAXTRAN`.

**Table 8-7 Characteristics of the DMTLOGDEV, DMTLOGNAME, DMTLOGSIZE, MAXRDTRAN, and MAXTRAN Parameters**

Parameter	Characteristics
DMTLOGDEV	Specifies the BEA Tuxedo file system that contains the Domains transaction log (DMTLOG) for this machine. The DMTLOG is stored as a BEA Tuxedo VTOC table on the device. If this parameter is not specified, the Domains gateway group is not allowed to process requests in transaction mode. Local domains running on the same machine can share the same DMTLOGDEV file system, but each local domain must have its own log (a table in the DMTLOGDEV) named as specified by the DMTLOGNAME keyword.
DMTLOGNAME	Specifies the name of the Domains transaction log for this domain. This name must be unique when the same DMTLOGDEV is used for several local domains. If a value is not specified, the value defaults to the string DMTLOG. The name must contain 30 characters or less.
DMTLOGSIZE	<p>Specifies the numeric size of the Domains transaction log for this machine (in pages). It must be greater than zero and less than the amount of available space on the BEA Tuxedo file system. If a value is not specified, the value defaults to 100 pages.</p> <p><b>Note:</b> The number of domains in a transaction determine the number of pages you must specify in the DMTLOGSIZE parameter. One transaction does not necessarily equal one log page.</p>
MAXRDTRAN	Specifies the maximum number of domains that can be involved in a transaction. It must be greater than zero and less than 32,768. If a value is not specified, the value defaults to 16.
MAXTRAN	Specifies the maximum number of simultaneous global transactions allowed on this local domain. It must be greater than or equal to zero, and less than or equal to the MAXGTT parameter specified in the TUXCONFIG file. If not specified, the default is the value of MAXGTT.

## Characteristics of the AUTOTRAN and TRANTIME Parameters (WebLogic Enterprise CORBA and Tuxedo Servers)

The `DM_REMOTE_SERVICES` section of the Domains configuration file identifies information on services *imported* and available on remote domains. Remote services are associated with a particular remote domain.

Table 8-8 describes the two transaction-related parameters in this section: `AUTOTRAN` and `TRANTIME`.

**Note:** For EJB and RMI applications, these settings are ignored.

**Table 8-8 Characteristics of the AUTOTRAN and TRANTIME Parameters**

Parameter	Characteristics
AUTOTRAN	Used by gateways to automatically start/terminate transactions for remote services. This capability is required if you want to enforce reliable network communication with remote services. You specify this capability by setting the <code>AUTOTRAN</code> parameter to <code>Y</code> in the corresponding remote service definition.
TRANTIME	Specifies the default timeout value in seconds for a transaction automatically started for the associated service. The value must be greater than or equal to zero, and less than 2147483648. The default is 30 seconds. A value of zero implies the maximum timeout value for the machine.

# Sample Distributed Application Using Transactions

This topic includes the following sections:

- RESOURCES Section
- MACHINES Section
- GROUPS and NETWORK Sections
- SERVERS, SERVICES, and ROUTING Sections

This topic describes a sample configuration file for the Bankapp application, a sample CORBA application that enables transactions and distributes the application over three sites. The application includes the following features:

- Data-dependent routing on ACCOUNT\_ID.
- Data distributed over three databases.
- BRIDGE processes communicating with the system via the ATMI interface.
- System administration from one site.

The configuration file includes seven sections: RESOURCES, MACHINES, GROUPS, NETWORK, SERVERS, SERVICES, and ROUTING.

**Note:** Although this sample is a CORBA application, the principles apply to EJB applications as well, except that the ROUTING section is *not* used in EJB applications, nor are the TRANTIME and AUTOTRAN parameters in the INTERFACES section.

## RESOURCES Section

The RESOURCES section shown in Listing 8-1 specifies the following parameters:

- MAXSERVERS, MAXSERVICES, and MAXGTT are less than the defaults. This makes the Bulletin Board smaller.

- MASTER is SITE3 and the backup master is SITE1.
- MODEL is set to MP and OPTIONS is set to LAN, MIGRATE. This allows a networked configuration with migration.
- BBLQUERY is set to 180 and SCANUNIT is set to 10. This means that DBBL checks of the remote BBLs are done every 1800 seconds (one half hour).

#### **Listing 8-1 Sample RESOURCES Section**

---

```
*RESOURCES
#
IPCKEY          99999
UID             1
GID             0
PERM            0660
MAXACCESSERS    25
MAXSERVERS      25
MAXSERVICES     40
MAXGTT          20
MASTER         SITE3, SITE1
SCANUNIT        10
SANTYSCAN       12
BBLQUERY        180
BLOCKTIME       30
DBBLWAIT        6
OPTIONS         LAN, MIGRATE
MODEL           MP
LDBAL           Y
```

---

## **MACHINES Section**

The MACHINES section shown in Listing 8-2 specifies the following parameters:

- TLOGDEVICE and TLOGNAME are specified, which indicate that transactions will be done.
- The TYPE parameters are all different, which indicates that encode/decode will be done on all messages sent between machines.

### **Listing 8-2 Sample MACHINES Section**

---

```
*MACHINES
Gisela      LMID=SITE1
            TUXDIR="/usr/tuxedo"
            APPDIR="/usr/home"
            ENVFILE="/usr/home/ENVFILE"
            TLOGDEVICE="/usr/home/TLOG"
            TLOGNAME=TLOG
            TUXCONFIG="/usr/home/tuxconfig"
            TYPE="3B600"

romeo      LMID=SITE2
            TUXDIR="/usr/tuxedo"
            APPDIR="/usr/home"
            ENVFILE="/usr/home/ENVFILE"
            TLOGDEVICE="/usr/home/TLOG"
            TLOGNAME=TLOG
            TUXCONFIG="/usr/home/tuxconfig"
            TYPE="SEQUENT"

juliet     LMID=SITE3
            TUXDIR="/usr/tuxedo"
            APPDIR="/usr/home"
            ENVFILE="/usr/home/ENVFILE"
            TLOGDEVICE="/usr/home/TLOG"
            TLOGNAME=TLOG
            TUXCONFIG="/usr/home/tuxconfig"
            TYPE="AMDAHL"
```

---

## **GROUPS and NETWORK Sections**

The GROUPS and NETWORK sections shown in Listing 8-3 specify the following parameters:

- The TMSCOUNT is set to 2, which means that only two TMS\_SQL transaction manager servers will be booted per group.
- The OPENINFO string indicates that the application will perform database access.



**Listing 8-3 Sample GROUPS and NETWORK Sections**

---

```
*GROUPS
DEFAULT:          TMSNAME=TMS_SQL          TMSCOUNT=2
BANKB1            LMID=SITE1                GRPNO=1
  OPENINFO="TUXEDO/SQL:/usr/home/bankd11:bankdb:readwrite"
BANKB2            LMID=SITE2                GRPNO=2
  OPENINFO="TUXEDO/SQL:/usr/home/bankd12:bankdb:readwrite"
BANKB3            LMID=SITE3                GRPNO=3
  OPENINFO="TUXEDO/SQL:/usr/home/bankd13:bankdb:readwrite"

*NETWORK
SITE1              NADDR="0X0002ab117B2D4359"
                  BRIDGE="/dev/tcp"
                  NLSADDR="0X0002ab127B2D4359"

SITE2              NADDR="0X0002ab117B2D4360"
                  BRIDGE="/dev/tcp"
                  NLSADDR="0X0002ab127B2D4360"

SITE3              NADDR="0X0002ab117B2D4361"
                  BRIDGE="/dev/tcp"
                  NLSADDR="0X0002ab127B2D4361"
```

---

## SERVERS, SERVICES, and ROUTING Sections

The `SERVERS`, `SERVICES`, and `ROUTING` sections shown in Listing 8-4 specify the following parameters:

- The TLR servers have a `-T` number passed to their `tpsrvrinit()` functions.
- All requests for the services are routed on the `ACCOUNT_ID` field.
- None of the services will be performed in `AUTOTRAN` mode.

**Note:** The `ROUTING` section is *not* used in EJB or RMI applications.

### **Listing 8-4 Sample SERVERS, SERVICES, and ROUTING Sections**

---

```
*SERVERS
DEFAULT: RESTART=Y MAXGEN=5 REPLYQ=N CLOPT="-A"
TLR      SRVGRP=BANKB1      SRVID=1      CLOPT="-A -- -T 100"
TLR      SRVGRP=BANKB2      SRVID=3      CLOPT="-A -- -T 400"
TLR      SRVGRP=BANKB3      SRVID=4      CLOPT="-A -- -T 700"
XFER     SRVGRP=BANKB1      SRVID=5      REPLYQ=Y
XFER     SRVGRP=BANKB2      SRVID=6      REPLYQ=Y
XFER     SRVGRP=BANKB3      SRVID=7      REPLYQ=Y

*SERVICES
DEFAULT: AUTOTRAN=N
WITHDRAW      ROUTING=ACCOUNT_ID
DEPOSIT       ROUTING=ACCOUNT_ID
TRANSFER      ROUTING=ACCOUNT_ID
INQUIRY       ROUTING=ACCOUNT_ID

*ROUTING
ACCOUNT_ID    FIELD=ACCOUNT_ID      BUFTYPE="FML"
              RANGES="MON - 9999:*,
              10000 - 39999:BANKB1
              40000 - 69999:BANKB2
              70000 - 100000:BANKB3
              ""
```

---

---

# Index

## A

- ACID properties 1-2, 2-3
- activation policies
  - transaction 3-11
- always transaction policy 3-4, 3-12
- API models, supported 1-3
- Application-to-Transaction Monitor
  - Interface (ATMI) 2-6
- assigning transaction policies 3-7
- atomicity (ACID properties) 1-2
- autocommit 7-6
- AUTOTRAN parameter 8-9, 8-10, 8-13

## B

- bean-managed transactions 1-10
  - transaction attributes 5-5
  - transaction semantics
    - stateful session beans 5-9
    - stateless session beans 5-10
- business transactions, support 1-4

## C

- client applications
  - multithreading 2-7
  - using transactions 4-5
- CLOPT parameter 7-13
- close\_xa\_rm method 3-10
- CLOSEINFO parameter 8-7
- closing an XA resource manager 3-10

- CMTRET parameter 8-3

- code example

- C++ CORBA server object that supports transactions 1-20

- EJB applications 1-25

- OMG IDL for Transactions sample CORBA application 1-16

- RMI applications 1-27

- TransactionCurrent object 4-3

- transactions

- C++ 4-6

- Java 4-6

- Visual Basic 4-7

- transactions in C++ CORBA client application 1-22

- UBBCONFIG file for Transactions sample CORBA application 1-23

- committing transactions

- EJB applications 1-27

- RMI applications 1-29

- connection pool

- finding via JNDI 7-16

- consistency (ACID properties) 1-2

- container-managed transactions 1-9

- transaction attributes 5-4

- transaction semantics 5-6

- entity beans 5-8

- stateful session beans 5-6

- stateless session beans 5-7

- Control interface 2-17

---

## CORBA applications

- transaction propagation 2-4
- Transaction Service 2-8
- Transaction Service API
  - Control interface 2-17
  - Current interface 2-12
  - data types 2-10
  - exceptions 2-10
  - TransactionalObject interface 2-18
- transactions in client applications 1-7, 4-2

- transactions overview 1-6
- Transactions sample CORBA application 1-13
- transactions support 3-2

## CORBAServices Object Transaction Service (OTS) 1-3, 2-2, 4-2

## CosTransactions module 2-12, 2-18

## Current interface

- about the Current interface 2-12
- begin method 2-13
- commit method 2-14
- get\_control method 2-15
- get\_status method 2-14
- get\_transaction\_name method 2-15
- resume method 2-17
- rollback method 2-14
- rollback\_only method 2-14
- set\_timeout method 2-15
- suspend method 2-16

customer support contact information xi

## D

### data types

CORBA Transaction Service API 2-10

### database cursors 3-4

deactivate\_object method  
and transactions 3-11

defining user-defined exceptions 3-13

delegated commit 2-3

## development process

### client applications

Transactions sample CORBA application 1-21

transactions 4-2

## distributed transactions 7-4

- about distributed transactions 1-4
- about implementing 7-14
- finding the connection pool via JNDI 7-16
- importing packages 7-15
- initializing TransactionCurrent 7-16
- performing 7-17
- setting up 7-17

DMTLOGDEV parameter 8-12

DMTLOGNAME parameter 8-12

DMTLOGSIZE parameter 8-12

documentation, where to find it x

domain transaction log, creating 8-5

DR\_TRANS\_ABORT 3-11

DR\_TRANS\_COMMITTING 3-11

durability (ACID properties) 1-2

## E

## EJB applications

- bean-managed transactions 1-10
- committing transactions 1-27
- container-managed transactions 1-9
- exceptions 5-11
- general guidelines 5-2
- importing packages 1-25
- JNDI lookup 1-26
- participating in a transaction 5-5
- rolling back transactions 1-27
- sample code 1-25
- session synchronization 5-10
- starting transactions 1-26
- timeouts 5-11
- transaction attributes 5-3
- transaction semantics 5-5

---

- transactions overview 1-8
- ENABLEXA parameter 7-5
- entity beans
  - container-managed transactions
    - transaction semantics 5-8
  - JDBC/XA accessibility 7-11
- exceptions
  - CORBA Transaction Service API 2-10
  - EJB applications 5-11
  - HeuristicHazard 2-11
  - HeuristicMixed 2-11
  - INVALID\_TRANSACTION 2-20
  - InvalidControl 2-12
  - NoTransaction 2-11
  - OBJ\_ADAPTER 2-20
  - SubtransactionsUnavailable 2-11
  - TRANSACTION\_ROLLEDBACK 2-16
  - Unavailable 2-12
  - user-defined exceptions 3-13
- explicit propagation, in CORBA applications 2-4

## F

- FACTORYROUTING parameter 8-9
- flat transactions 2-5

## G

- global transactions 7-4
- GROUPS section 8-16

## H

- handling exceptions
  - EJB applications 5-11
- HeuristicHazard exception 2-11
- HeuristicMixed exception 2-11

## I

- ICF file
  - defining transaction policies 4-2
- ignore transaction policy 3-7
- Implementation Configuration File (ICF)
  - defining transaction policies 1-18
- implicit propagation, in CORBA applications 2-4
- importing packages
  - EJB applications 1-25
- interdomain interoperability 2-5
- interoperability
  - interdomain 2-5
  - intradomain 2-5
  - network 2-6
  - remote clients and WLE domain 2-5
- intradomain interoperability 2-5
- INVALID\_TRANSACTION exception 2-20
- InvalidControl exception 2-12
- isolation (ACID properties) 1-2

## J

- Java Naming Directory Interface (JNDI)
  - EJB applications 1-26
  - RMI applications 1-29
- Java Transaction API (JTA) 1-3, 2-2
- Java Transaction Service (JTS) 1-3
- JavaServerXA 7-12
- JavaServerXA, described 7-3
- JDBC standards, supported 7-4
- JDBC/XA driver
  - enabling 7-5
  - pooled connections 7-3
- JNDI
  - finding the connection pool 7-16

## L

- lightweight clients
  - about lightweight clients 2-3

---

- Application-to-Transaction Monitor
  - Interface (ATMI) 2-6
- interoperability 2-5
- listings
  - sample GROUPS section 8-17
  - sample MACHINES section 8-16
  - sample NETWORK section 8-17
  - sample RESOURCES section 8-15
- local transactions 7-4

## M

- MACHINES section 8-4, 8-15
- Mandatory transaction attribute 5-4
- MAXGTT parameter 8-3
- MAXRDTRAN parameter 8-12
- MAXTRAN parameter 8-12
- multithreading
  - clients 2-7

## N

- nested transactions 2-5, 2-20
- network interoperability 2-6
- NETWORK section 8-16
- Never transaction attribute 5-4
- never transaction policy 3-6
- NoTransaction exception 2-11
- NotSupported transaction attribute 5-4
- NULL resource manager 3-11

## O

- OBJ\_ADAPTER exception 2-20
- object state management 3-10
  - delegating to an XA RM 3-10
  - Transactions University sample application 3-17
- ODMG standard 2-7
- OMG IDL
  - Transactions sample CORBA

- application 1-15
- Open Group XA interface 2-6
- open\_xa\_rm method 3-8
- OPENINFO parameter 7-12, 8-7
- opening an XA resource manager 3-8
- optional transaction policy 3-5, 3-12
- Oracle7 3-18
- OSI TP protocol 2-6

## P

- participating in a transaction 5-5
- pooled connections 7-3
- printing product documentation x
- process failure, handling 2-7
- programming models, supported 1-2

## R

- recursive transactions 2-20
- related information xi
- remote clients and interoperability 2-5
- Required transaction attribute 5-4
- RequiresNew transaction attribute 5-4
- resource manager
  - closing an XA 3-10
  - delegating object state management 3-10
  - NULL 3-11
  - opening XA 3-8
- RESOURCES section 8-14
- RMI applications
  - committing transactions 1-29
  - general guidelines 6-2
  - JNDI lookup 1-29
  - rolling back transactions 1-29
  - sample code 1-27
  - starting transactions 1-29
  - transactions overview 1-11
- rolling back transactions
  - EJB applications 1-27
  - RMI applications 1-29

---

ROUTING parameter 8-10

ROUTING section 8-17

## S

sample applications

- Transactions sample CORBA application 1-13

Server Description File

- about the Server Description File 1-19

Server object

- supporting databases 1-19

- Transactions sample CORBA application 1-19

SERVERS section 8-17

SERVICES section 8-17

session synchronization 5-10

setTransactionTimeout method 5-11

SNA LU 6.2 protocol 2-7

SRVTYPE parameter 7-13

starting transactions

- EJB applications 1-26

- RMI applications 1-29

stateful session beans

- bean-managed transactions transaction semantics 5-9

- container-managed transactions transaction semantics 5-6

- JDBC/XA accessibility 7-10

stateless session beans

- bean-managed transactions transaction semantics 5-10

- container-managed transactions transaction semantics 5-7

- JDBC/XA accessibility 7-11

SubtransactionsUnavailable exception 2-11

support

technical xi

Supported transaction attribute 5-4

## T

terminating transactions 2-4

throwing user-defined exceptions 3-14

TLOG 8-3

TLOGDEVICE parameter 1-23, 8-5

TLOGNAME parameter 8-5

TLOGOFFSET parameter 8-5

TLOGSIZE parameter 8-5

TMS 3-18

- configuring 3-8

- Oracle7 3-8

- requirements for 3-8

TMSCOUNT parameter 8-7

TMSNAME parameter 8-7

transaction activation policy 3-11

transaction attributes

- bean-managed transactions 5-5

- container-managed transactions 5-4 described 5-3

transaction log, creating 8-3

Transaction Manager Server

- See TMS

transaction policies

- always 3-4, 3-12

- assigning 3-7

- defined 1-18

- defining in ICF file 4-2

- ignore 3-7

- never 3-6

- optional 3-5, 3-12

- Transactions sample CORBA application 1-18

transaction semantics 5-5

Transaction Service

- about the Transaction Service 2-2

- capabilities 2-2

- CORBA applications 2-8

---

- API extensions 2-19
- clients supported 2-8
- features 1-4
- general constraints 2-7
- limitations 2-2
- transactional objects
  - defining 3-4
- TransactionalObject interface 2-18
- TransactionCurrent interface 7-16
- TransactionCurrent object
  - begin method 4-4
  - code examples 4-3
  - commit method 4-4
  - get\_status method 4-5
  - get\_transaction\_name method 4-5
  - getcontrol method 4-5
  - getting initial references to 2-9
  - methods 4-4
  - resume method 4-5
  - rollback method 4-4
  - rollback\_only method 4-4
  - set\_timeout method 4-5
  - suspend method 4-5
- transaction-related parameters in
  - MACHINES section, defining 8-4
- transactions
  - autocommit 7-6
  - client CORBA applications 4-2
  - configuring
    - AUTOTRAN parameter 8-9, 8-10, 8-13
    - CMTRET parameter 8-3
  - creating a transaction log
    - creating the domain transaction log 8-5
    - creating the Universal Device List (UDL) 8-4
    - defining transaction-related parameters in MACHINES section 8-4
  - defining each resource manager and

- the transaction manager
  - server in GROUPS section 8-6
- DMTLOGDEV parameter 8-12
- DMTLOGNAME parameter 8-12
- DMTLOGSIZE parameter 8-12
- enabling a Tuxedo service to begin a transaction in the
  - SERVICES section 8-9
- FACTORYROUTING parameter 8-9
- INTERFACES section 8-8
- MAXGTT parameter 8-3
- MAXRDTRAN parameter 8-12
- MAXTRAN parameter 8-12
- modifying the domain configuration file to support transactions 8-11
- modifying the UBBCONFIG file 8-2
- ROUTING parameter 8-10
- sample GROUPS section 8-6
- specifying application-wide transactions in
  - RESOURCES 8-3
- TLOGDEVICE parameter 8-5
- TLOGNAME parameter 8-5
- TLOGOFFSET parameter 8-5
- TLOGSIZE parameter 8-5
- transaction log (TLOG) 8-3
- transaction values description in
  - sample GROUPS section 8-6
- TRANTIME parameter 8-9, 8-10, 8-13

- CORBA applications 1-6
- distributed
  - sample application 8-14
- distributed transactions 7-4
- EJB applications 1-8
- flat transactions 2-5



---

- functional overview 1-6
- implementing in a WLE server
  - application
- in client applications 4-5
- in CORBA client applications 1-7
- integrity 2-4
- local transactions 7-4
- nested 2-20
- nested transactions 2-5
- object state management 3-10
- participating in a transaction 5-5
- propagating, in CORBA applications 2-4
- recursive 2-20
- RMI applications 1-11
- termination 2-4
- timeouts 5-11
- transaction contexts, in JDBC/XA
  - connections 7-7
- transaction processing 2-6
- transaction semantics 5-5
- when to use transactions 1-5
- Transactions CORBA sample application
  - workflow 1-13
- Transactions sample CORBA application
  - about the Transactions sample CORBA application 1-13
  - client application 1-21
  - development steps 1-15
  - illustrated 1-13
  - location 1-15
  - OMG IDL 1-15
  - transaction policies 1-18
  - UBBCONFIG file 1-22
  - user exceptions 1-13
  - writing server applications 1-19
- Transactions University sample application
  - about the application 3-15
  - configuration requirements 3-19
  - object state management 3-17
  - transaction model used 3-16
- transactions, configuring

- CLOSEINFO parameter 8-7
- OPENINFO parameter 8-7
- TMSCOUNT parameter 8-7
- TMSNAME parameter 8-7
- trans-timeout-seconds element 5-11
- TRANTIME parameter 8-9, 8-10, 8-13
- two-phase commit protocol (2PC) 1-4

## U

- UBBCONFIG file
  - adding transactions 1-22
  - JDBCConnPOOLS section 7-5
- UDL 8-4
- Unavailable exception 2-12
- Universal Device List (UDL) 8-4
- unmanaged desktops 2-3
- user exceptions
  - Transactions sample CORBA application 1-13
- user-defined exceptions
  - about user-defined exceptions 3-13
  - defining 3-13
  - throwing 3-14
- UserTransaction
  - committing transactions
    - EJB applications 1-27
    - RMI applications 1-29
  - initializing
    - EJB applications 1-26
  - rolling back transactions
    - EJB applications 1-27
    - RMI applications 1-29
  - sample code 1-25, 1-27
  - starting transactions
    - EJB applications 1-26
    - RMI applications 1-29

## V

- vetoing a transaction 3-11

---

## W

### WLE JDBC/XA driver

- about the driver 7-2

- accessibility

  - CORBA methods 7-9

  - EJB methods 7-9

- using 7-12

### WLE server applications

- and transactions

## X

### XA resource manager

- closing 3-10

- delegating object state management 3-10

- opening 3-8

- Transactions University sample

  - application 3-18